

CS231A

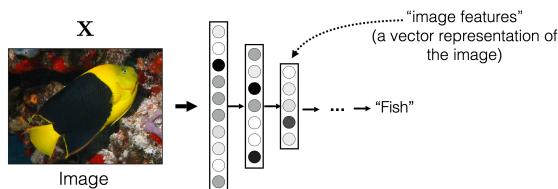
Computer Vision: From 3D Reconstruction to Recognition



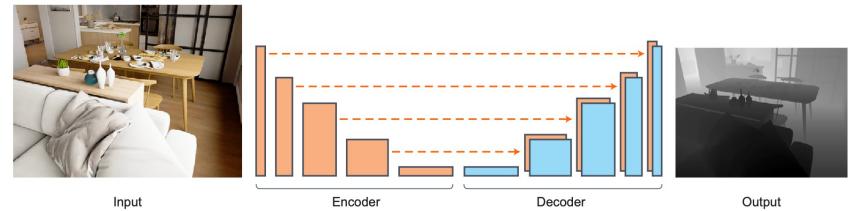
Optical and Scene Flow

Learning Goals for Upcoming Lectures

Representations & Representation Learning



Monocular Depth Estimation, Feature Tracking

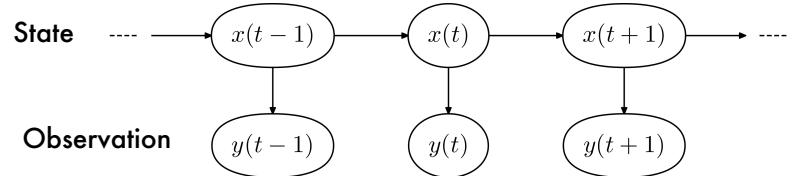


Optical & Scene Flow

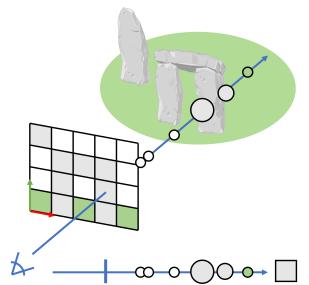


A Database and Evaluation Methodology for Optical Flow.
Baker et al. IJCV. 2011

Optimal Estimation



Neural Radiance Fields



What will you learn today?

Optical Flow

What is it and why do you care?

Assumptions

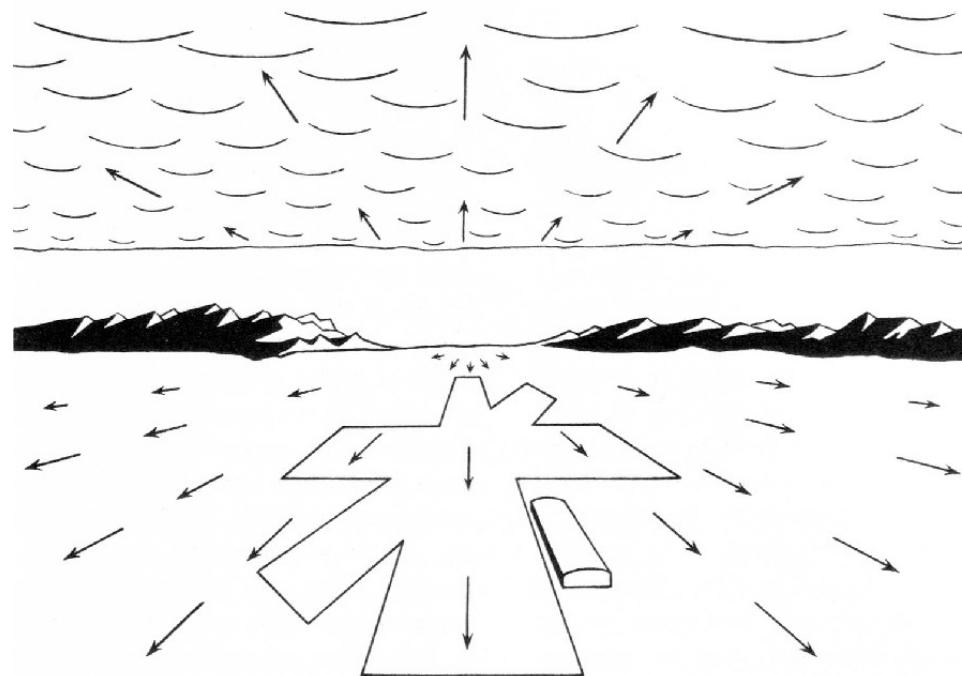
Formulating the optimization problem

Solving it

Scene Flow

Learning-based Approaches to Estimating Motion

Optical Flow - What is it?



J. J. Gibson, The Ecological Approach to Visual Perception

Optical Flow - What is it?

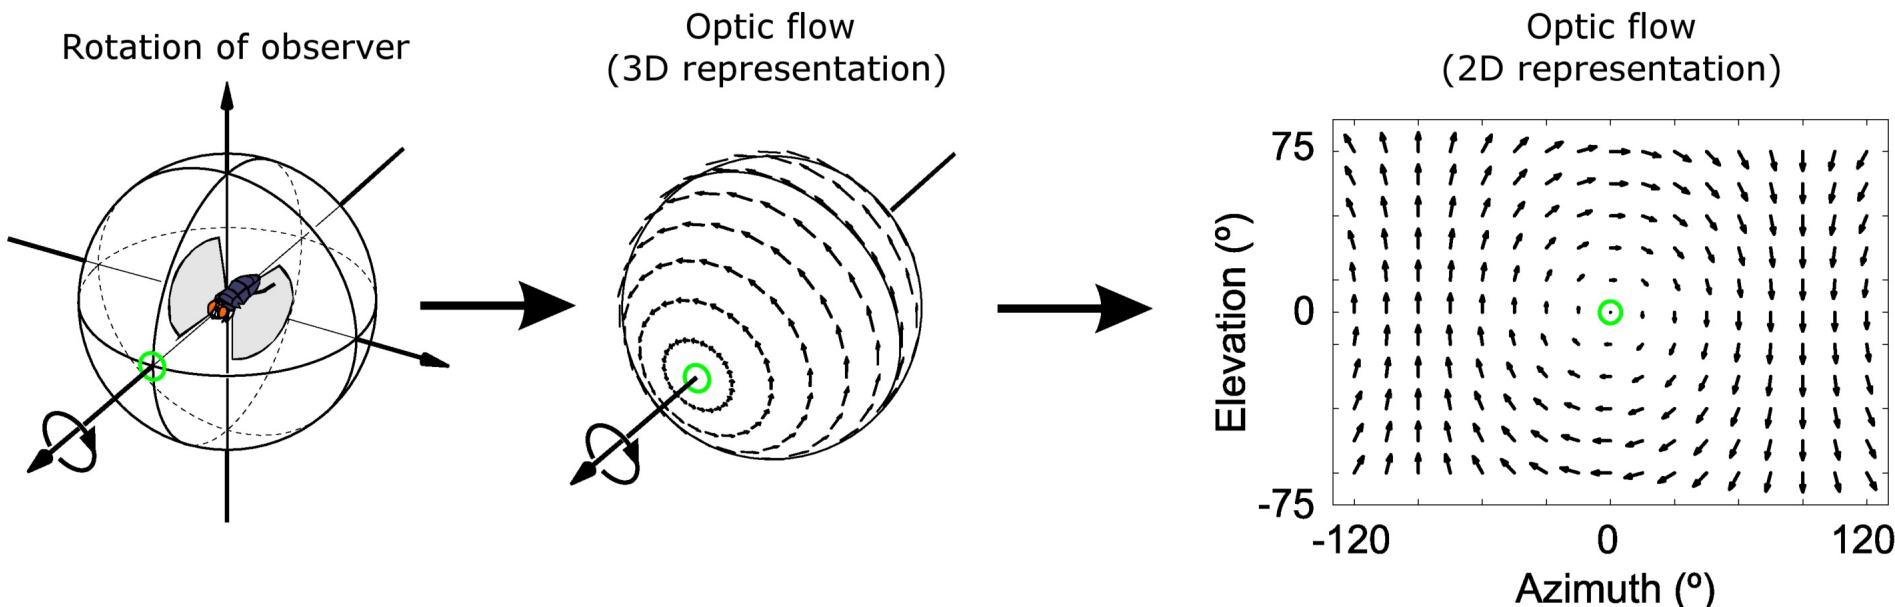
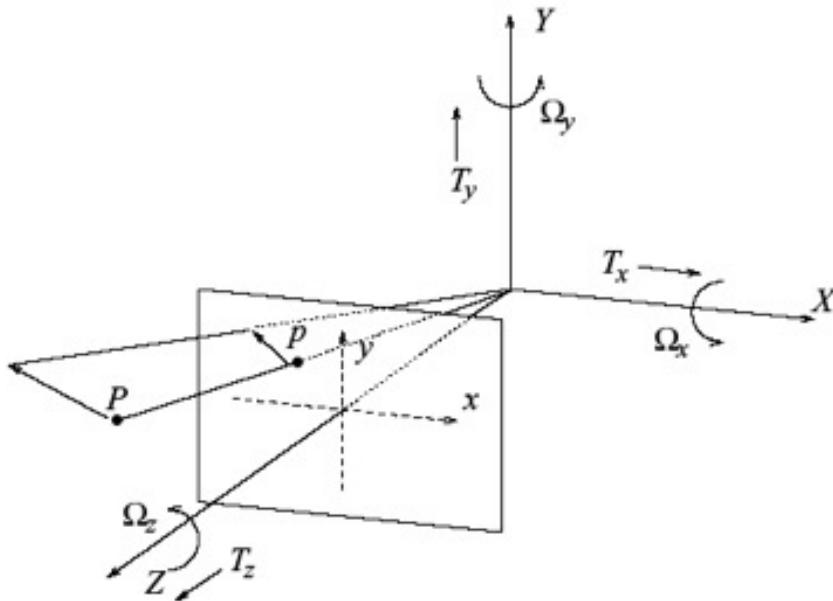


Image Credit: Wikipedia. Optical Flow.

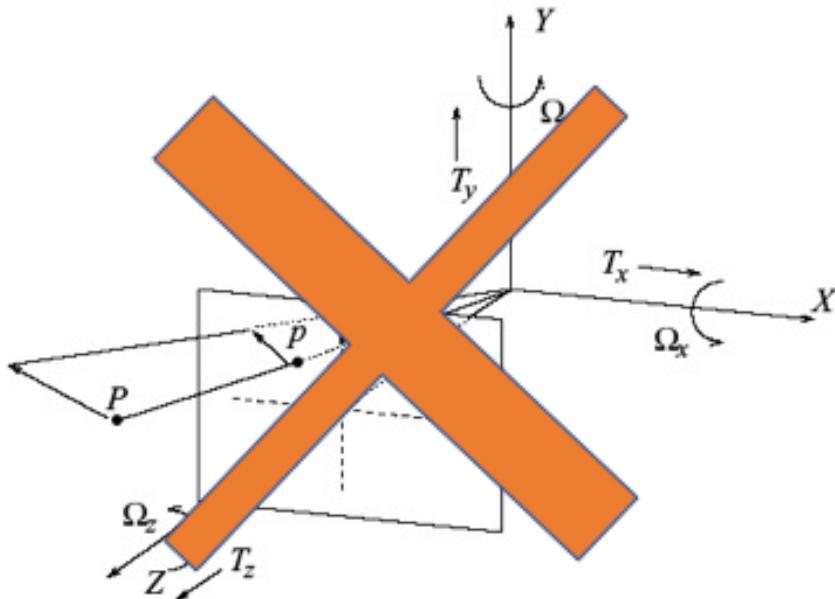
Motion Field



Motion field = 2D motion field representing the projection of the 3D motion of points in the scene onto the image plane.

B. Horn, Robot Vision, MIT Press

Optical flow



Optical flow = 2D velocity field describing the **apparent** motion in the images.

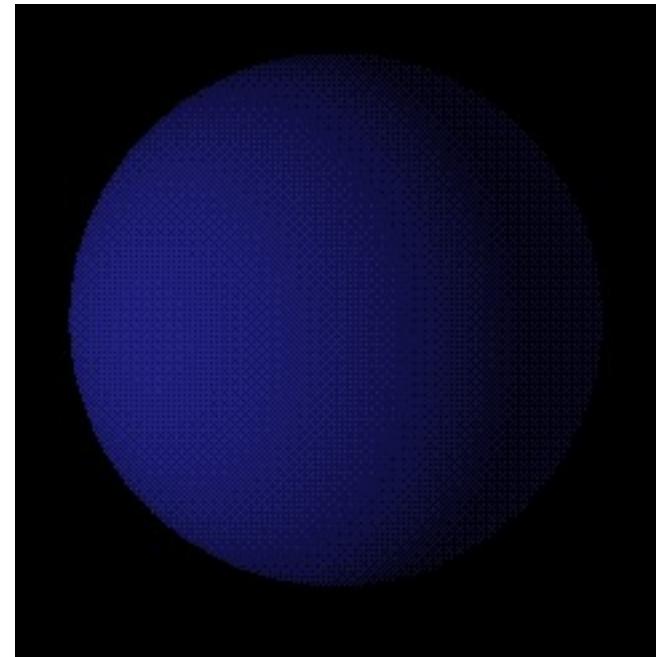
B. Horn, Robot Vision, MIT Press

What is the motion field? What is the apparent motion?

Lambertian (matte) ball rotating in 3D

What does the 2D motion field look like?

What does the 2D optical flow field look like?



Slide Credit: Michael Black

Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

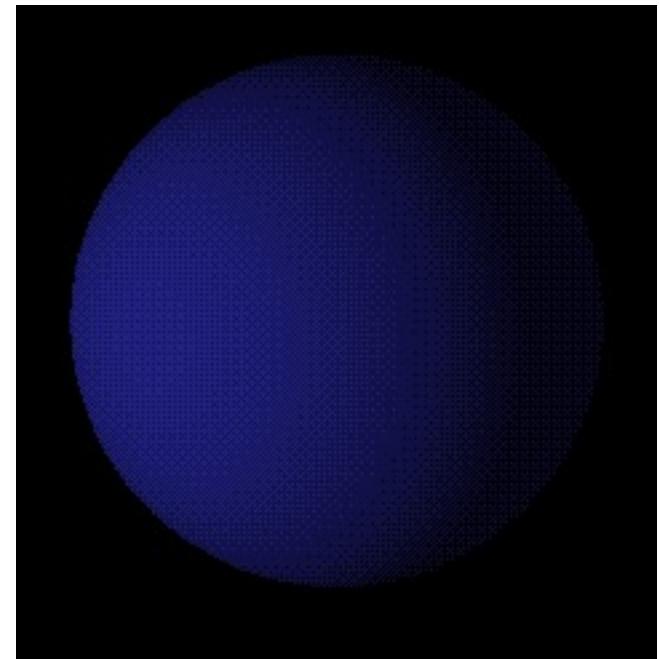
What is the motion field? What is the apparent motion?

Stationary Lambertian (matte) ball

Moving Light Source

What does the 2D motion field look like?

What does the 2D optical flow field look like?



Slide Credit: Michael Black

Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

Optical flow - What is it?

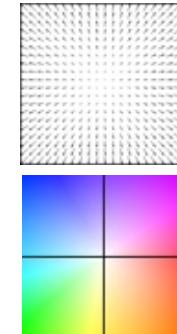
Motion Displacement of all image pixels



Image pixel value at time t and
Location $\mathbf{x} = (x, y)$: $I(x, y, t)$



$u(x, y)$ horizontal component
 $v(x, y)$ vertical component

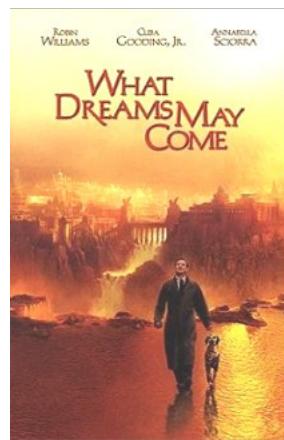


Key

Slide Credit: Michael Black

Optical Flow - What is it good for?

Painterly effect



Slide Credit: Michael Black

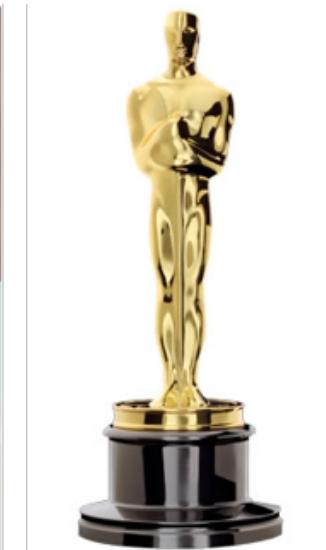
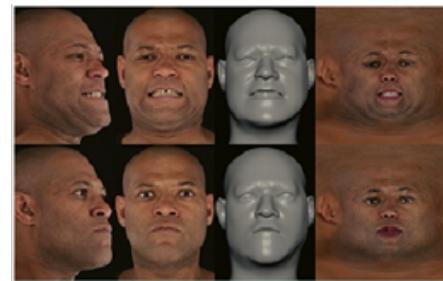
Optical Flow - What is it good for?

Face morphing in matrix reloaded



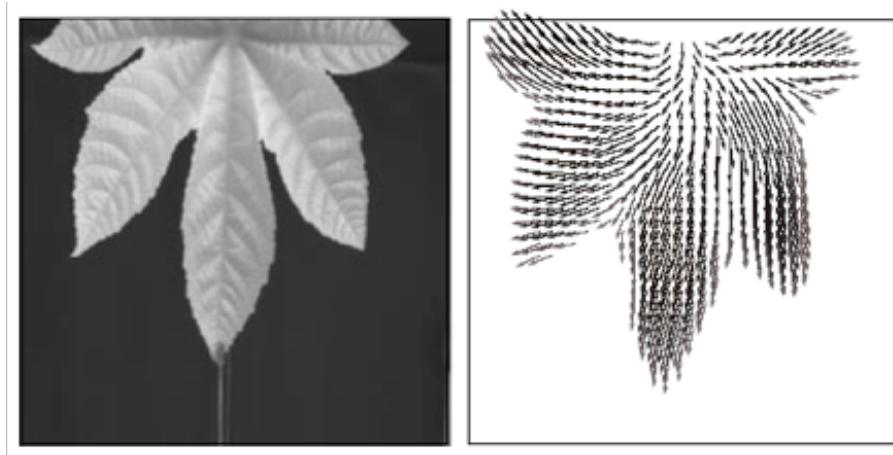
George Borshukov, Dan Piponi, Oystein Larsen, J.P.Lewis, Christina Tempelaar-Lietz
ESC Entertainment

SIGGRAPH'03



Slide Credit: Michael Black

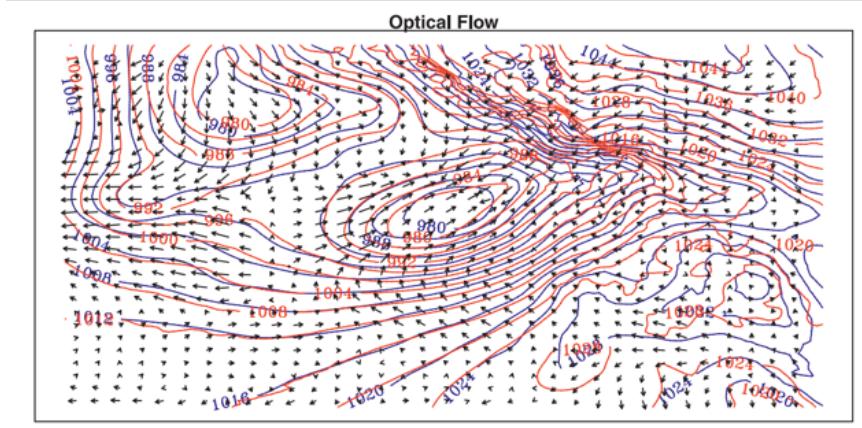
Optical Flow - What is it good for?



Tensor-based Image Sequence Processing Techniques for the Study of
Dynamical Processes, Haussecker, Spies, and Jahne, Proc. Intern. Symp.
On Real-time Imaging and Dynamic Analysis: 704-711, 1998

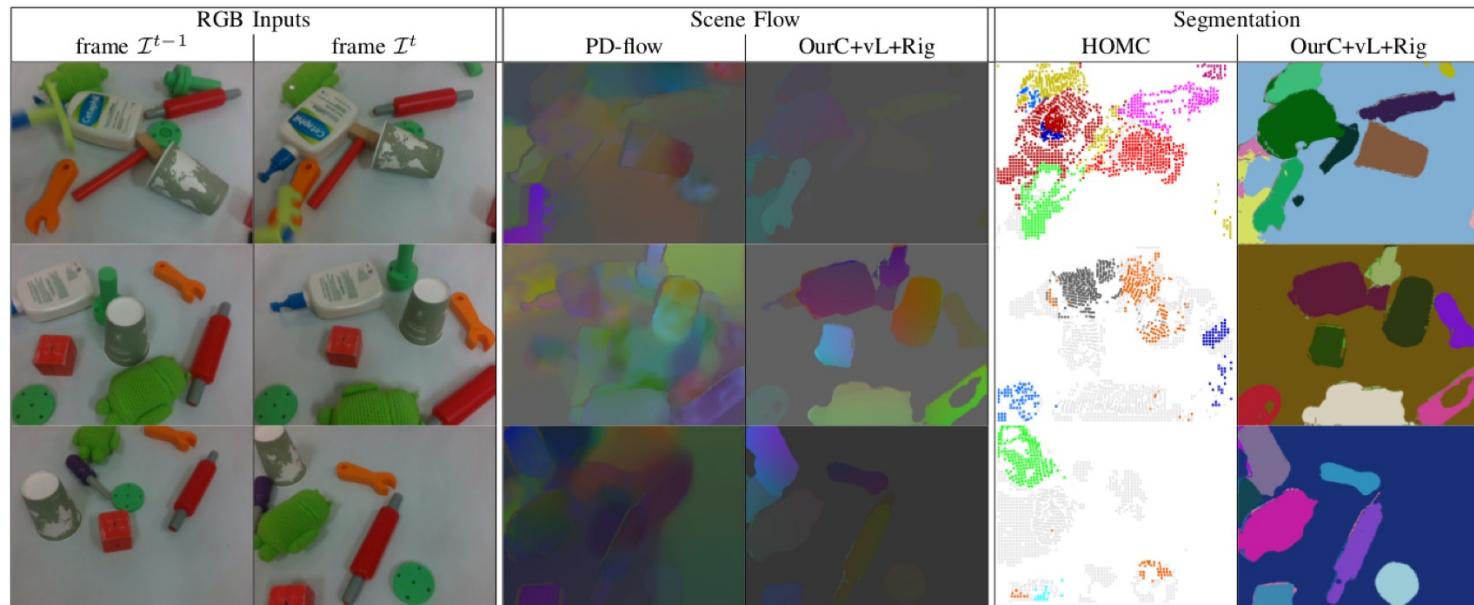
Slide Credit: Michael Black

Optical Flow - What is it good for?

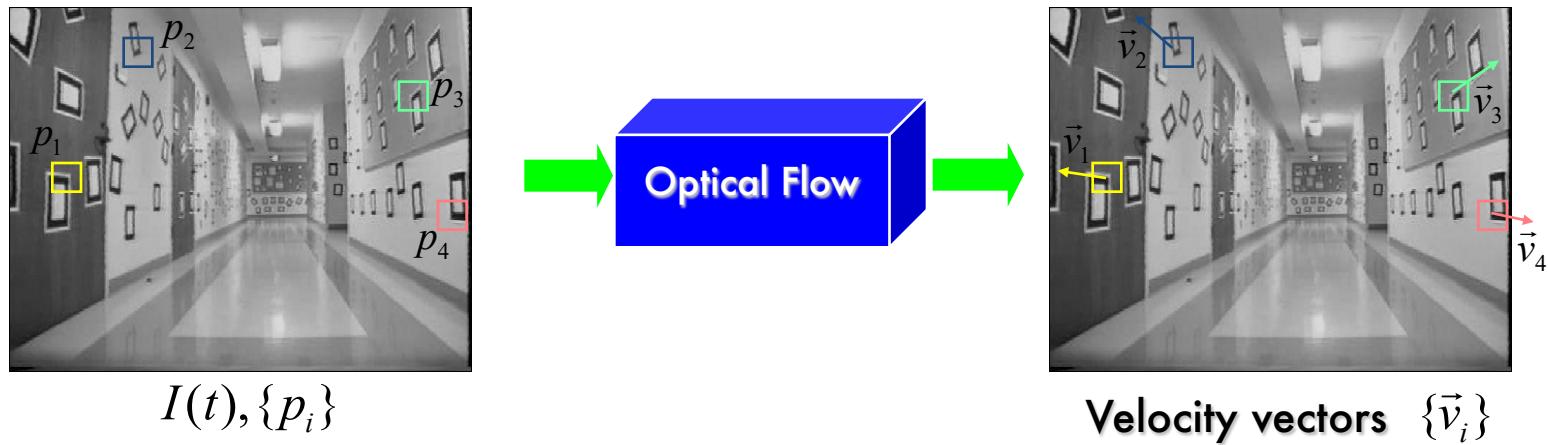


Slide Credit: Michael Black

Optical Flow - What is it good for?



Optical Flow - What is it good for?



Slide Credit: CS223b – Sebastian Thrun

Compute Optical Flow

Goal

Compute the **apparent** 2D image motion of pixels from one image frame to the next in a video sequence.

Compute (Sparse) Optical Flow

Also see CS131a

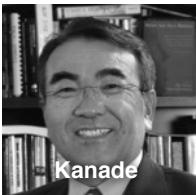


$I(t), \{p_i\}$



Velocity vectors $\{\vec{v}_i\}$

Simple KLT Tracker



An Iterative Image Registration Technique
with an Application to Stereo Vision.

1981



Detection and Tracking of Feature Points.

1991

The original KLT algorithm



Good Features to Track.

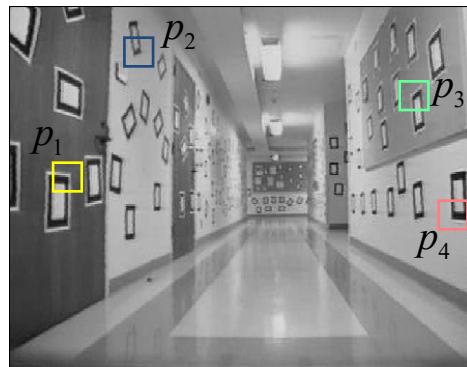
1994

16-385 Computer Vision (Kris Kitani)

Simple KLT Tracker

1. Find good points to track (Harris corners)
2. For each Harris corner compute motion (translation or affine) between consecutive frames
3. Link motion vector of successive frames to get a track for each Harris point
4. Introduce new Harris points by running detector every 10-15 frames
5. Track old and new corners using step 1-3

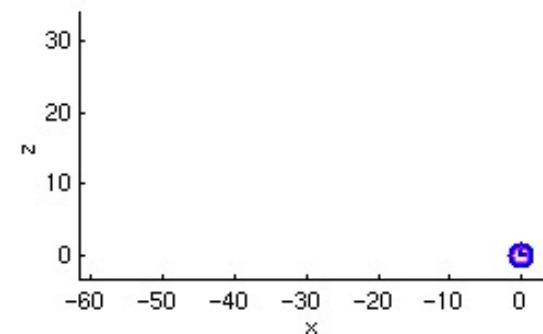
Computing (Sparse) Optical Flow



$I(t), \{p_i\}$



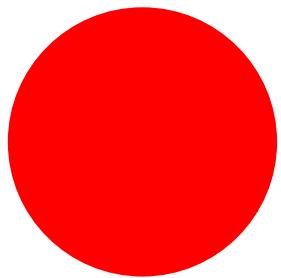
Velocity vectors $\{\vec{v}_i\}$



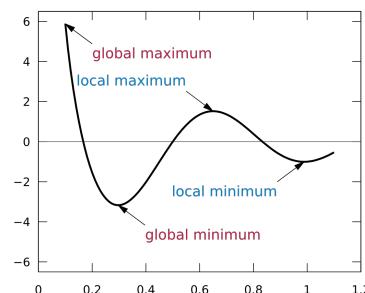
Jean-Yves Bouguet, [Ph.D.](#) CalTech

Compute (Dense) Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function

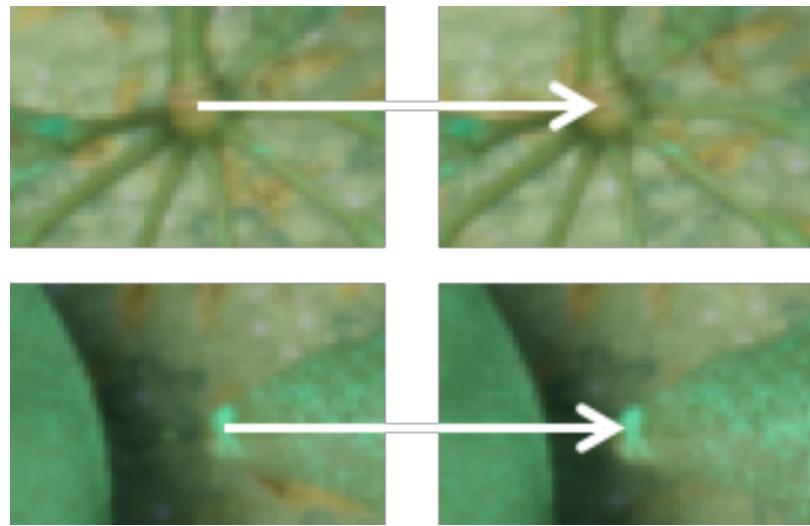


Source: Wikipedia.

Step 3 - Optimization



Assumption 1 - Brightness Constancy



$$I(x + u, y + v, t + 1) = I(x, y, t)$$

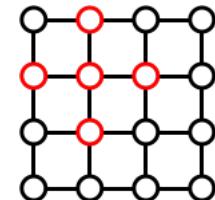
u, v = pixel offset t = time x, y = pixel position

Slide Credit: Michael Black

Assumption 2 - Spatial Smoothness



- Neighboring pixels in the image are likely to belong to the same surface.
- Surfaces are mostly smooth.
- Neighboring pixels will have similar flow.



Slide Credit: Michael Black

Assumption 3 – Temporal Coherence

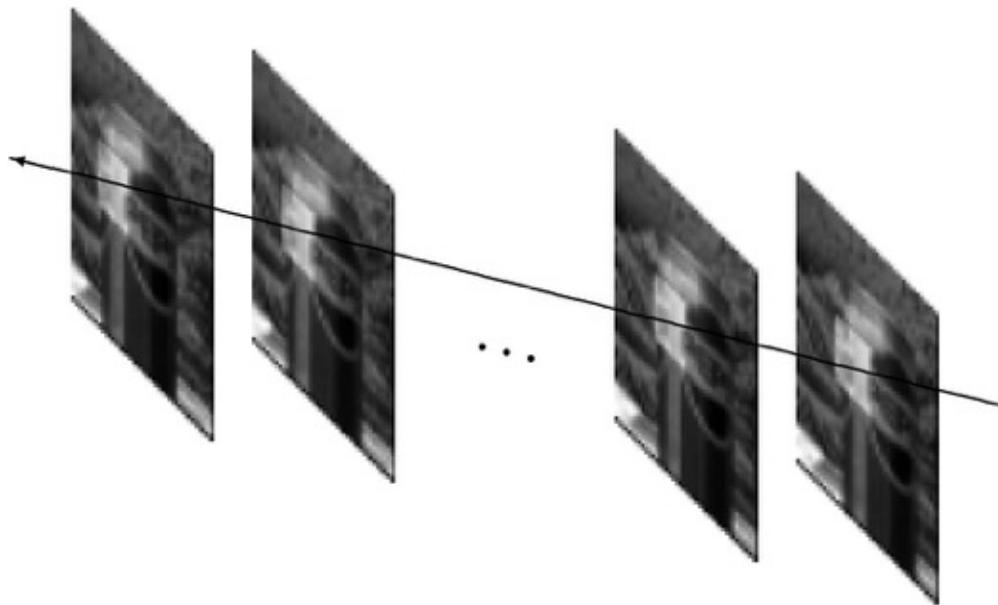
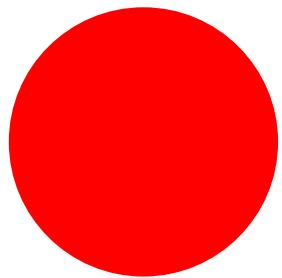


Figure 1.8: Temporal continuity assumption. A patch in the image is assumed to have the same motion (constant velocity, or acceleration) over time.

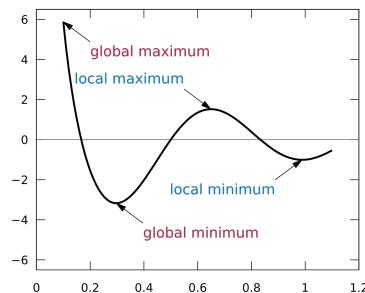
Slide Credit: Michael Black

Compute Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function



Source: Wikipedia.

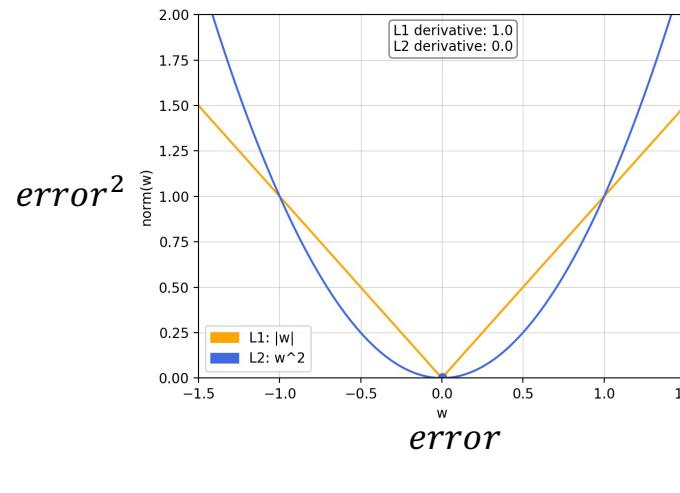
Objective Function – Data term - Brightness Constancy

\mathbf{u}, \mathbf{v} = optical flow field

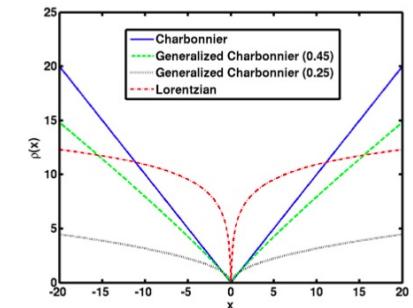
$$E_D(\mathbf{u}, \mathbf{v}) = \sum_{S=\text{all pixels}} (I(x_s + u_s, y_s + v_s, t+1) - I(x, y, t))^2$$

New Assumption: Quadratic error implies Gaussian noise

Quadratic penalty



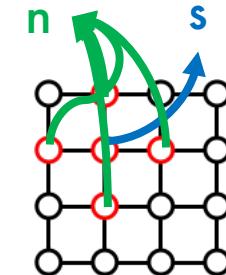
Alternative:
Huber/L1 Loss



Objective Function – Spatial Term – Spatial Smoothness

$$E_S(\mathbf{u}, \mathbf{v}) = \sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2$$

$G(s)$ = Pixel Neighborhood



New Assumptions:
Flow field smooth
Gaussian Deviations
First order smoothness good enough
Flow derivative approximated by first differences

Objective Function

$$E(u, v) = \underline{E_D(u, v)} + \lambda \underline{E_S(u, v)}$$

Optimization Variables
Relative weighting term

$$E(u, v) = \sum_s (I(x_s + u_s, y_s + v_s, t + 1) - I(x, y, t))^2 + \lambda (\sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2)$$

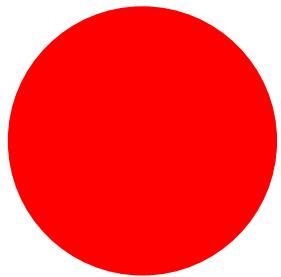
Data term

Spatial term

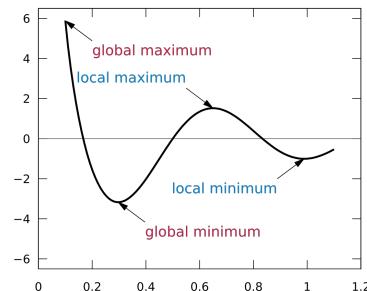
Nonlinear Optimization

Compute Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function



Source: Wikipedia.

Step 3 - Optimization



Linear Approximation

$$E(u, v) = E_D(u, v) + \lambda E_S(u, v)$$

$$E(u, v) = \sum_s (I(x_s + u_s, y_s + v_s, t + 1) - I(x, y, t))^2 + \lambda (\sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2)$$
$$u_s = dx, v_s = dy, dt = 1$$

Partial Derivative
in x direction

Partial Derivative
in y direction

$$\cancel{I(x, y, t) + dx \frac{\delta}{\delta x} I(x, y, t) + dy \frac{\delta}{\delta y} I(x, y, t) + dt \frac{\delta}{\delta t} I(x, y, t)} - I(x, y, t) = 0$$



Constraint Equation for Optical Flow

Optical Flow Constraint Equation

Linearized cost function

$$u \frac{\delta}{\delta x} I(x, y, t) + v \frac{\delta}{\delta y} I(x, y, t) + \frac{\delta}{\delta t} I(x, y, t) = 0$$

$I_x u + I_y v + I_t = 0$ = Constraint at every pixel

New Assumptions:

Flow is small

Image is differentiable

First order Taylor series is a good approximation

Optical Flow Constraint Equation

Notation

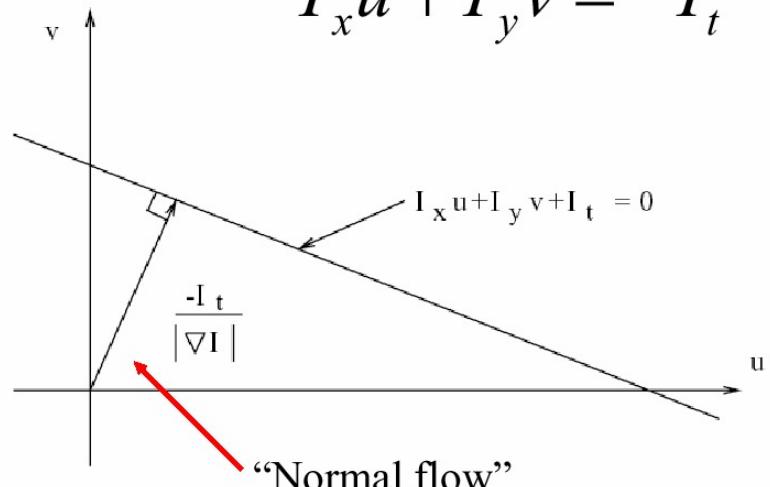
$$I_x u + I_y v + I_t = 0$$

$$\nabla I^T \mathbf{u} = -I_t$$

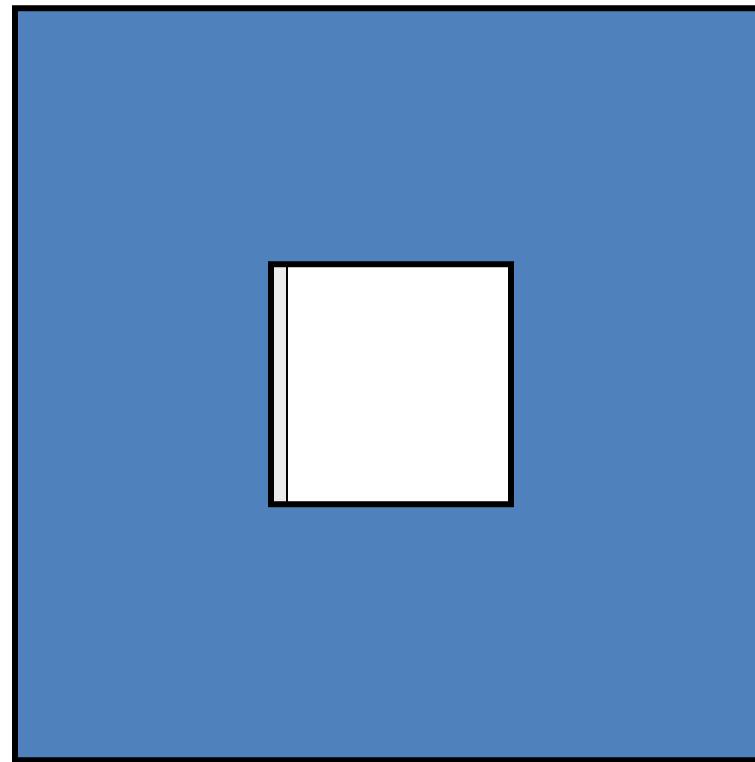
$$\mathbf{u} = \begin{bmatrix} u \\ v \end{bmatrix} \quad \nabla I = \begin{bmatrix} I_x \\ I_y \end{bmatrix}$$

At a single image pixel, we get a line:

$$I_x u + I_y v + I_t = 0$$

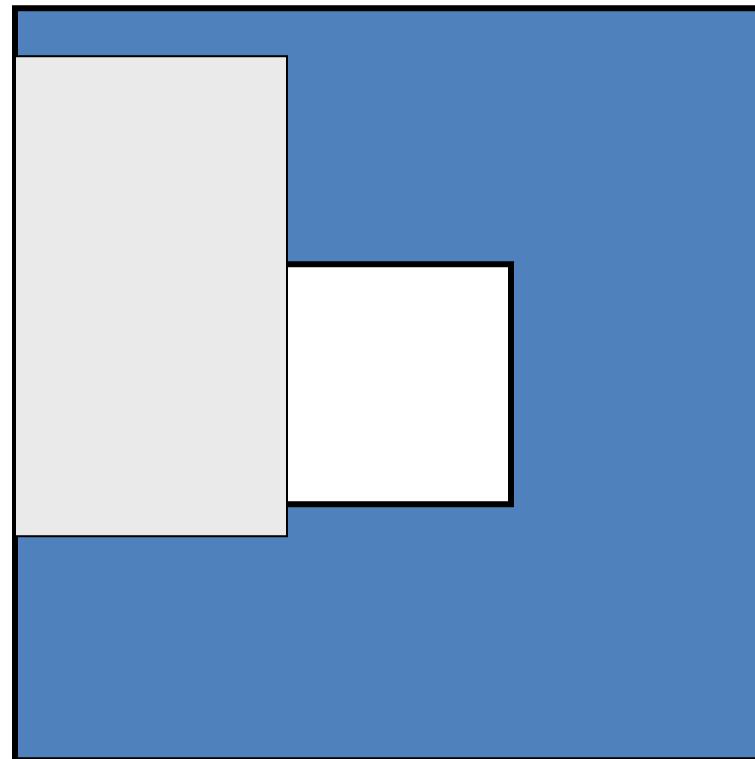


Aperture Problem



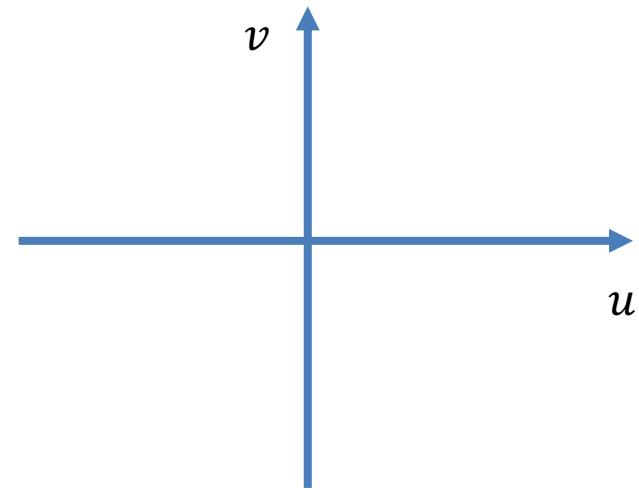
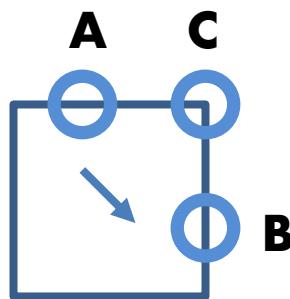
Slide Credit: CS223b – Sebastian Thrun

Aperture Problem

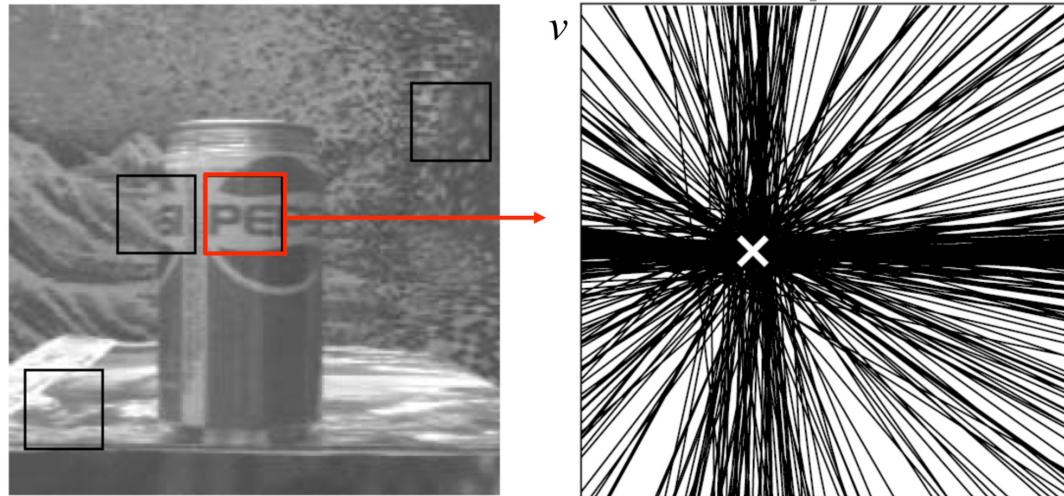


Slide Credit: CS223b – Sebastian Thrun

What are the constraint lines?



Multiple Constraints



Each pixel gives us a constraint: $I_x u + I_y v = -I_t$

Slide Credit: Michael Black

Image Gradient Examples - Edge

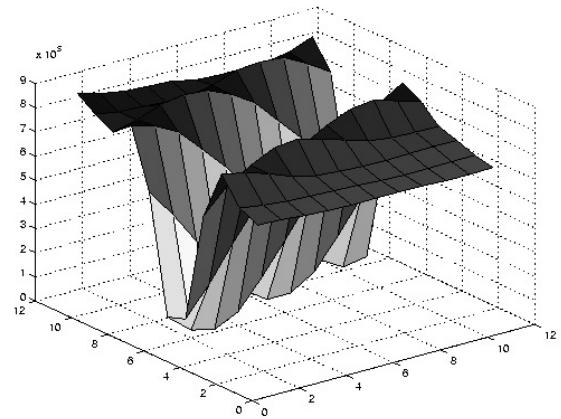
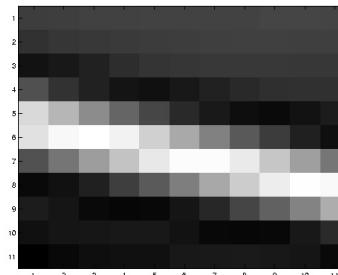


Image Gradient Examples – Low texture

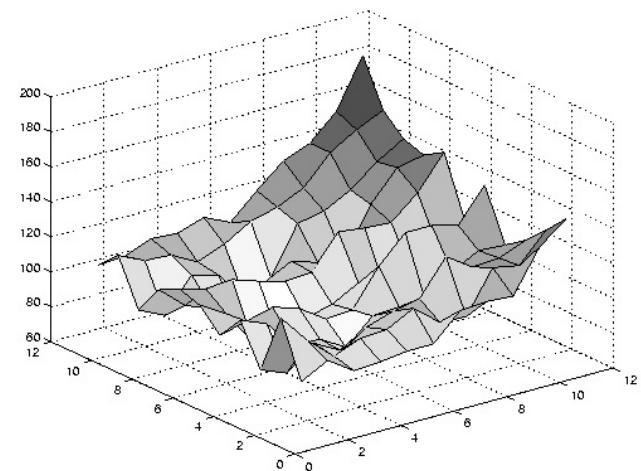
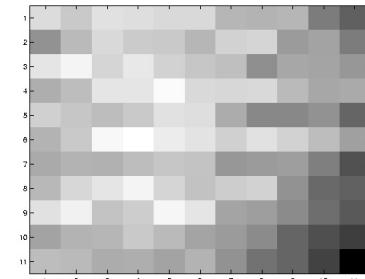
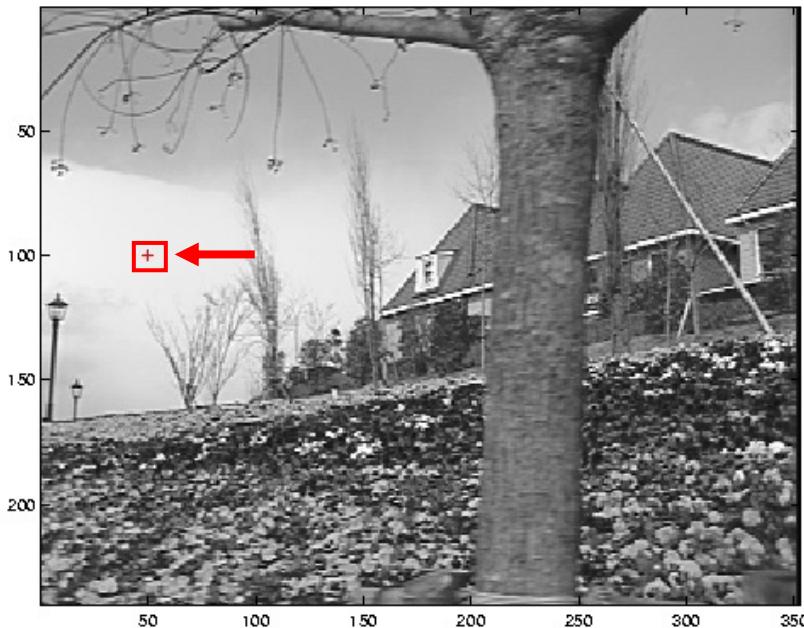
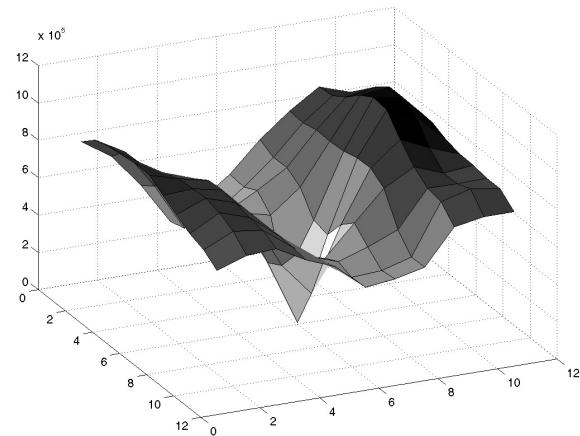
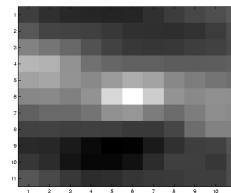


Image Gradient Examples – Low texture



How do we solve this optimization problem?

$$E(u, v) = \sum_{x, y \in R} (I_x(x, y, t)u + I_y(x, y, t)v + I_t(x, y, t))^2$$

$$\frac{\partial E}{\partial u} = \sum_R (I_x u + I_y v + I_t) I_x = 0$$

$$\frac{\partial E}{\partial v} = \sum_R (I_x u + I_y v + I_t) I_y = 0$$

Horn-Schunk Method

How do we solve this optimization problem?

Rearrange in Matrix form

$$\left[\sum_R I_x^2 \right] u + \left[\sum_R I_x I_y \right] v = - \sum_R I_x I_t$$

$$\left[\sum_R I_x I_y \right] u + \left[\sum_R I_y^2 \right] v = - \sum_R I_y I_t$$

$$\begin{bmatrix} \sum R I_x^2 & \sum R I_x I_y \\ \sum R I_y I_x & \sum R I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} - \sum R I_x I_t \\ - \sum R I_y I_t \end{bmatrix}$$

How do we solve this optimization problem?

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum I_x I_t \\ -\sum I_y I_t \end{bmatrix}$$

$$\mathbf{A}\mathbf{u} = \mathbf{b}$$

↑
unknown

How do we solve this optimization problem?

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum I_x I_t \\ -\sum I_y I_t \end{bmatrix}$$
$$\mathbf{A}\mathbf{u} = \mathbf{b}$$

If A was invertible

$$\mathbf{A}^{-1} \mathbf{A} \mathbf{u} = \mathbf{A}^{-1} \mathbf{b}$$

$$\mathbf{u} = \mathbf{A}^{-1} \mathbf{b} \qquad \qquad \mathbf{A} \mathbf{u} = \mathbf{b}$$

$$\mathbf{A}^T \mathbf{A} \mathbf{u} = \mathbf{A}^T \mathbf{b}$$

$$\mathbf{u} = \underline{(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}}$$

Pseudoinverse

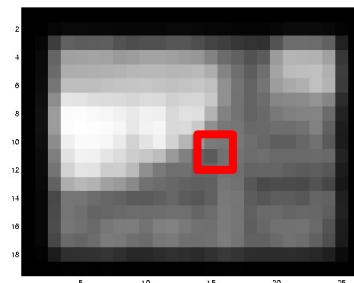
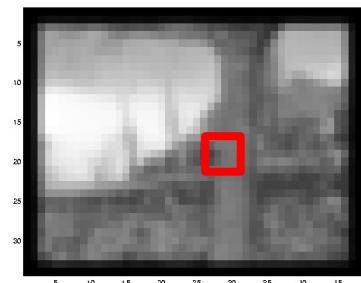
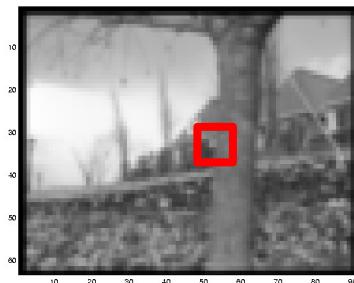
Bag of tricks

Small motion assumption



Bag of tricks

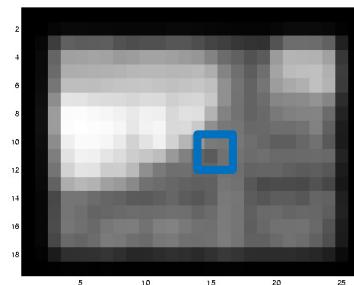
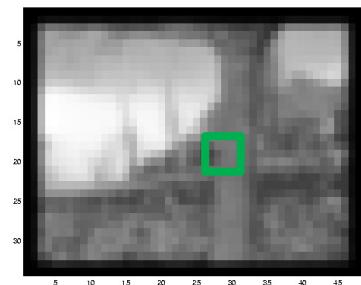
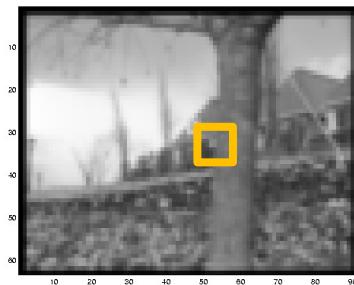
Reduce Resolution



* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Bag of tricks

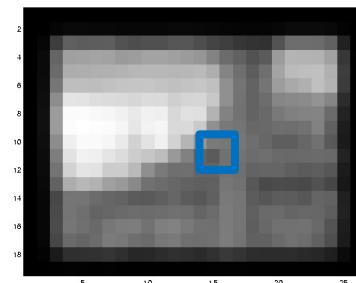
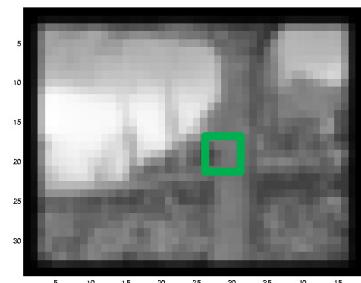
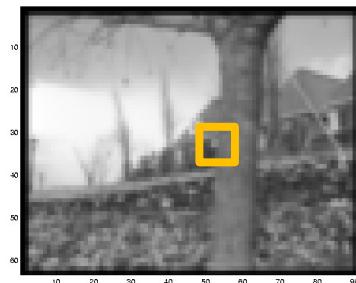
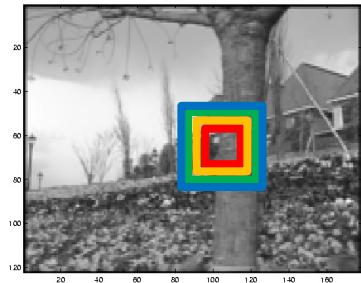
Reduce Resolution



* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

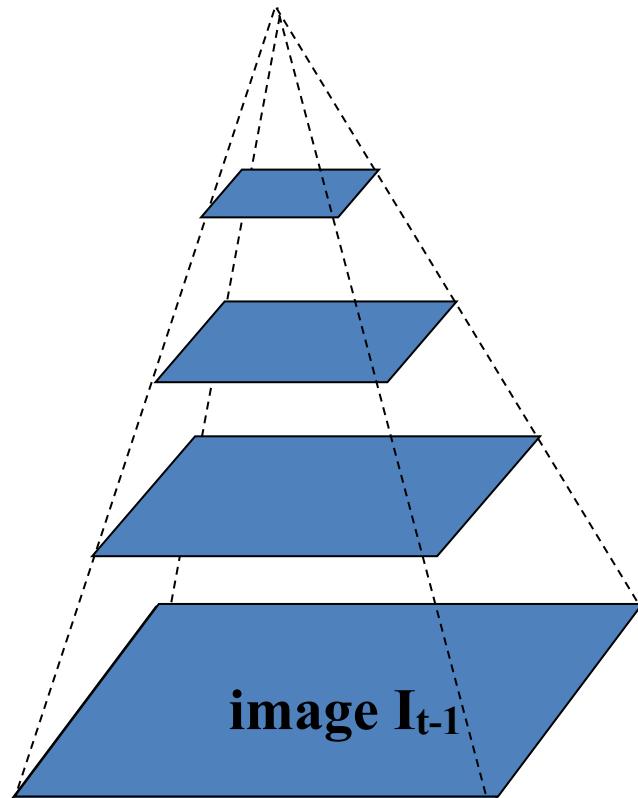
Bag of tricks

Reduce Resolution



* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Spatial Pyramids



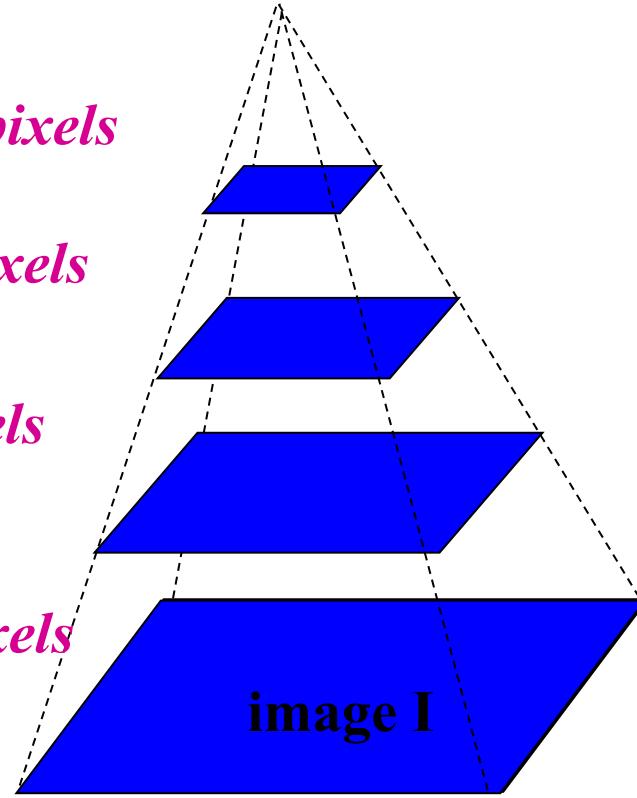
Gaussian pyramid of image I_{t-1}

$u=1.25 \text{ pixels}$

$u=2.5 \text{ pixels}$

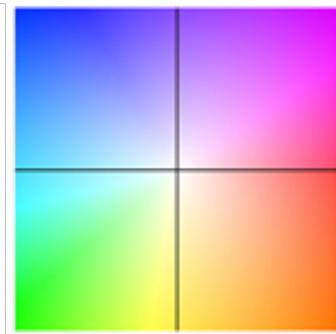
$u=5 \text{ pixels}$

$u=10 \text{ pixels}$

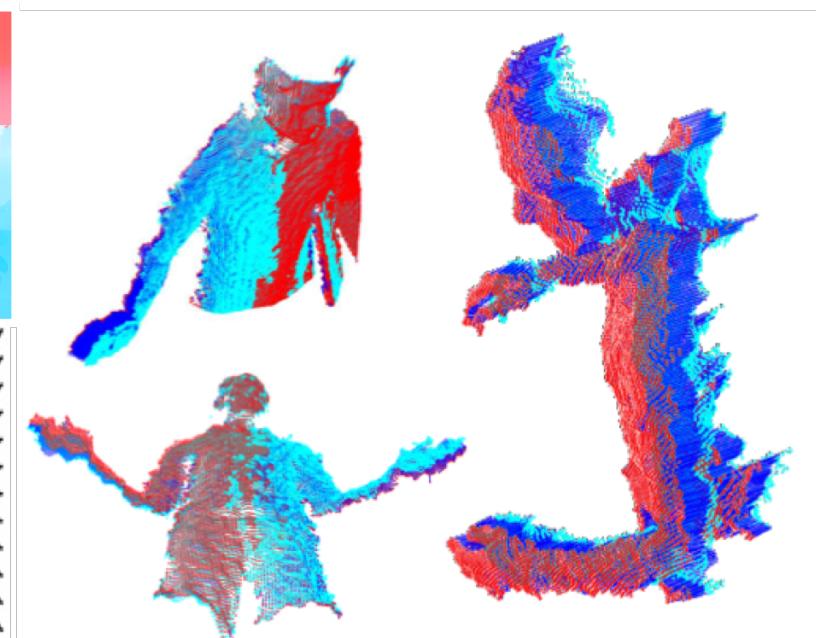


Gaussian pyramid of image I

Scene Flow = 3D Optical Flow



A Database and Evaluation Methodology for Optical Flow.
Baker et al. IJCV. 2011



A Primal-Dual Framework for Real-Time Dense RGB-D Scene Flow. Jaimez et al. ICRA, 2015.

What are the main challenges with this traditional formulation?

- Assumptions
 - Brightness constancy
 - Small motion
 - Etc
- Occlusions
- Large motion

Learning-based approaches

- Since 2015 - FlowNet
- Availability of synthetic data, e.g. Sintel

FlowNet - Learning Optical Flow with Convolutional Networks

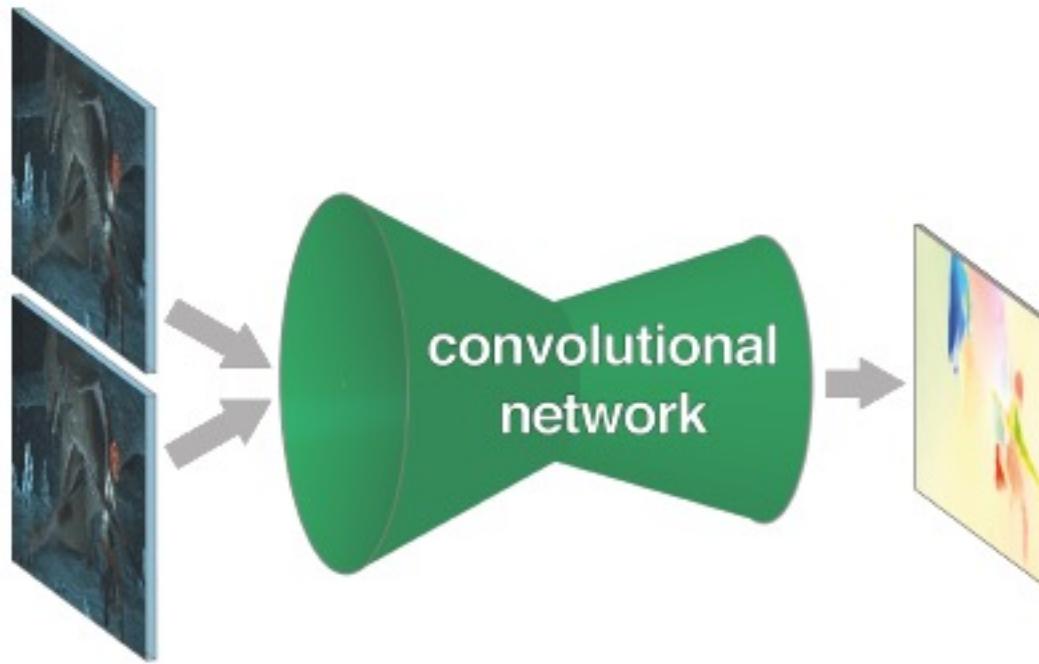


Image Pair

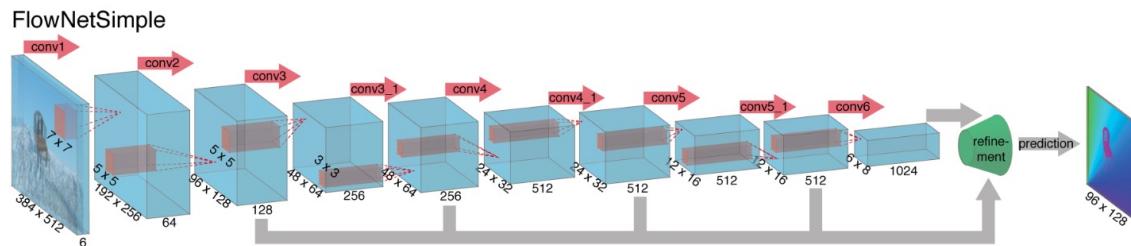
Optical Flow

Supervised Learning with Labeled Data Set

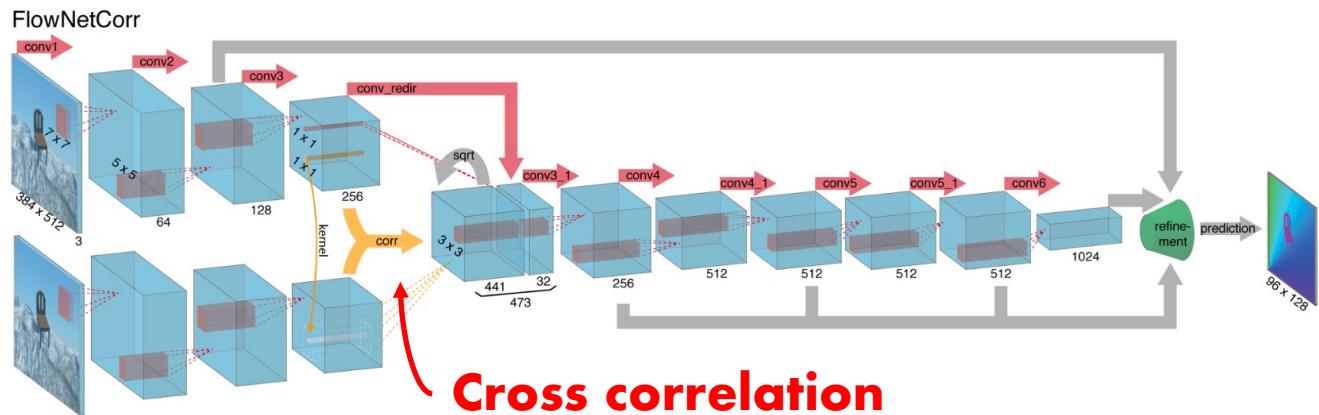
Alexey Dosovitskiy, Philipp Fischer, Eddy Ilg, P. Häusser, C. Hazırbaş, V. Golkov, P. Smagt, D. Cremers, Thomas Brox. IEEE International Conference on Computer Vision (ICCV), 2015

FlowNet - Learning Optical Flow with Convolutional Networks

1



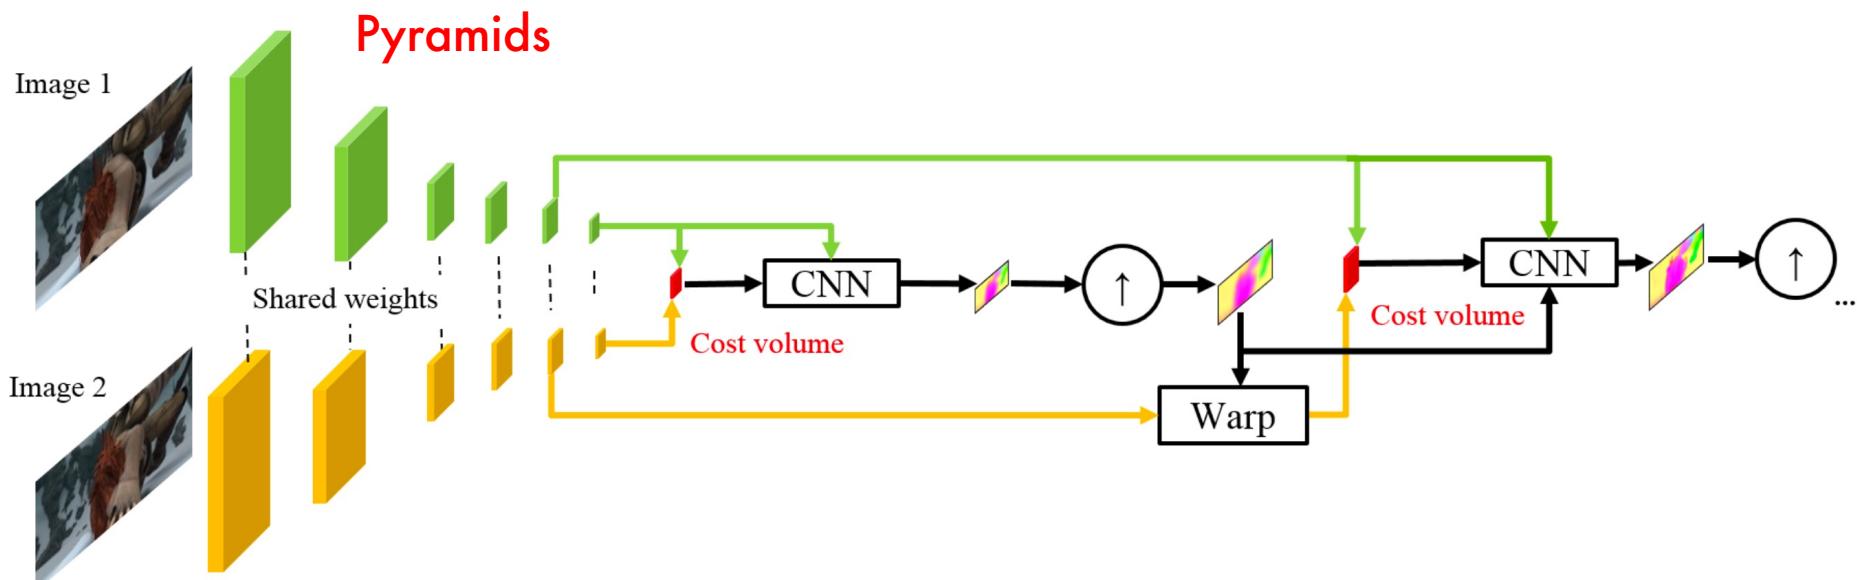
2



Supervised Learning

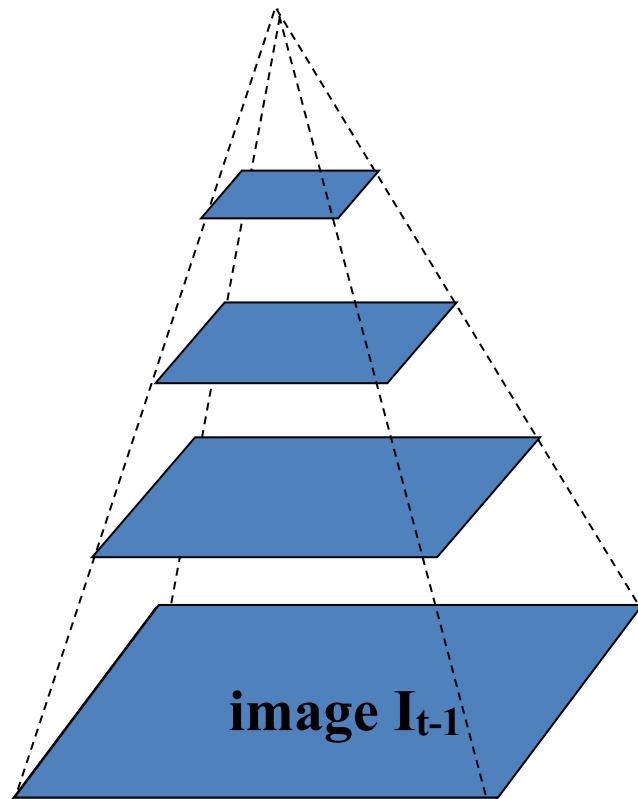
Supervised Optical Flow using traditional principles

Pyramidal Processing, Image Warping & Cost volume processing



Deqing Sun, Xiaodong Yang, Ming-Yu Liu, and Jan Kautz. "PWC-Net: CNNs for Optical Flow Using Pyramid, Warping, and Cost Volume." CVPR 2019

Spatial Pyramids



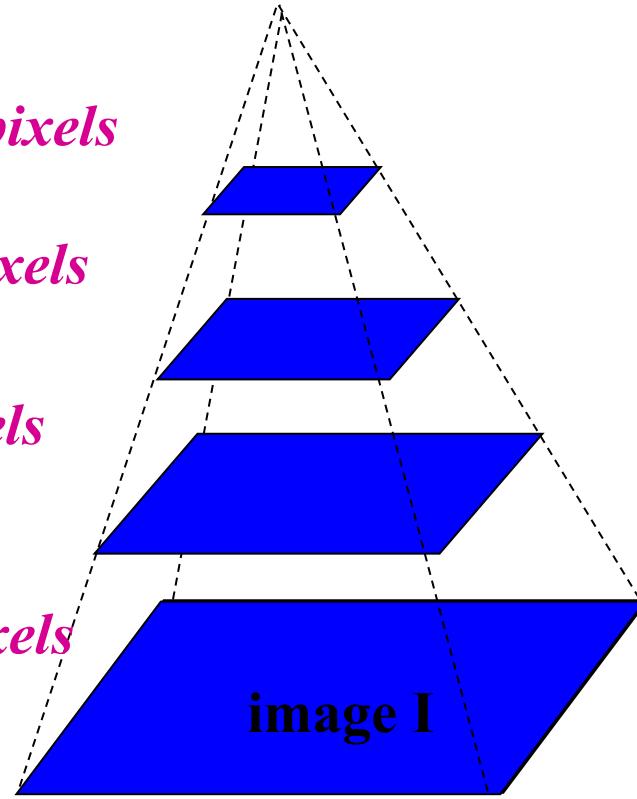
Gaussian pyramid of image I_{t-1}

$u=1.25 \text{ pixels}$

$u=2.5 \text{ pixels}$

$u=5 \text{ pixels}$

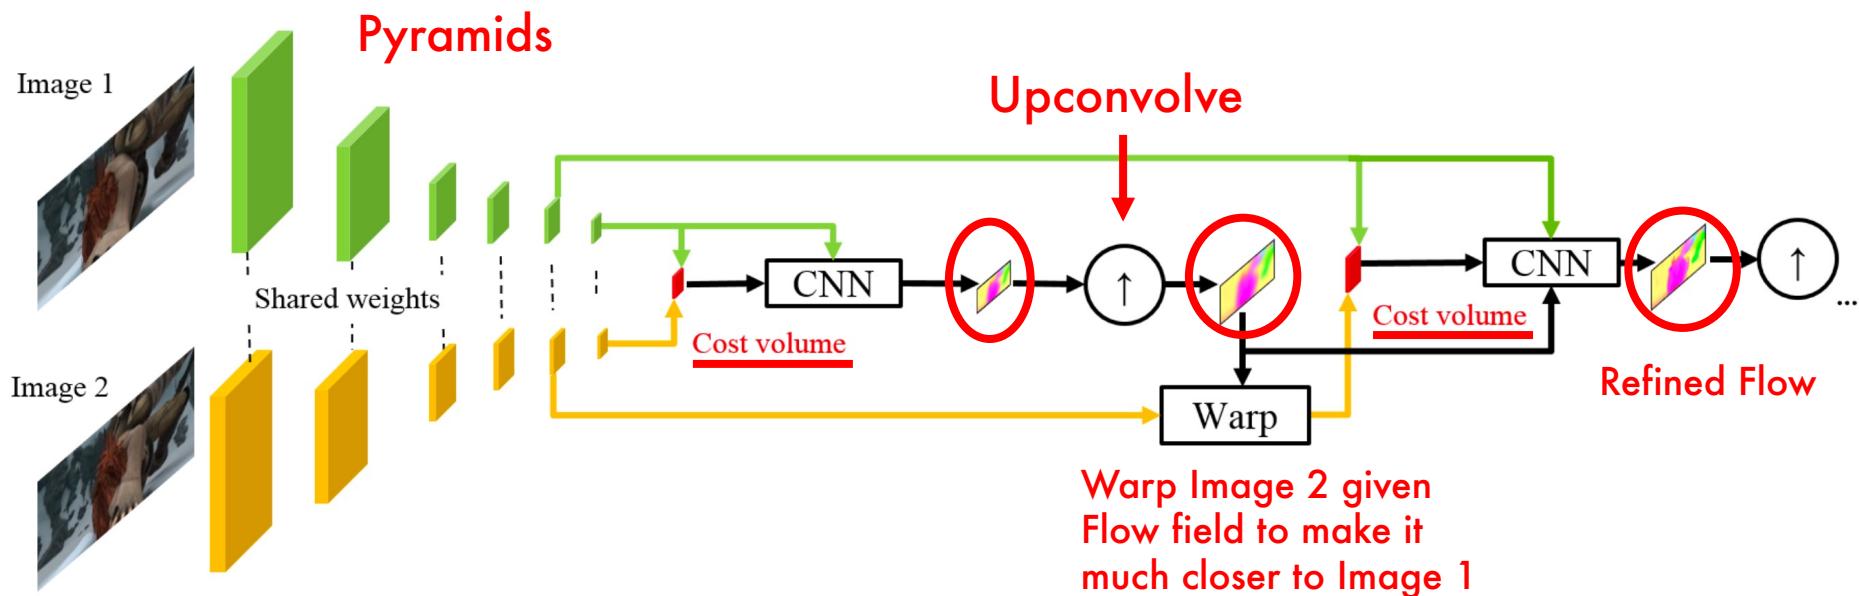
$u=10 \text{ pixels}$



Gaussian pyramid of image I

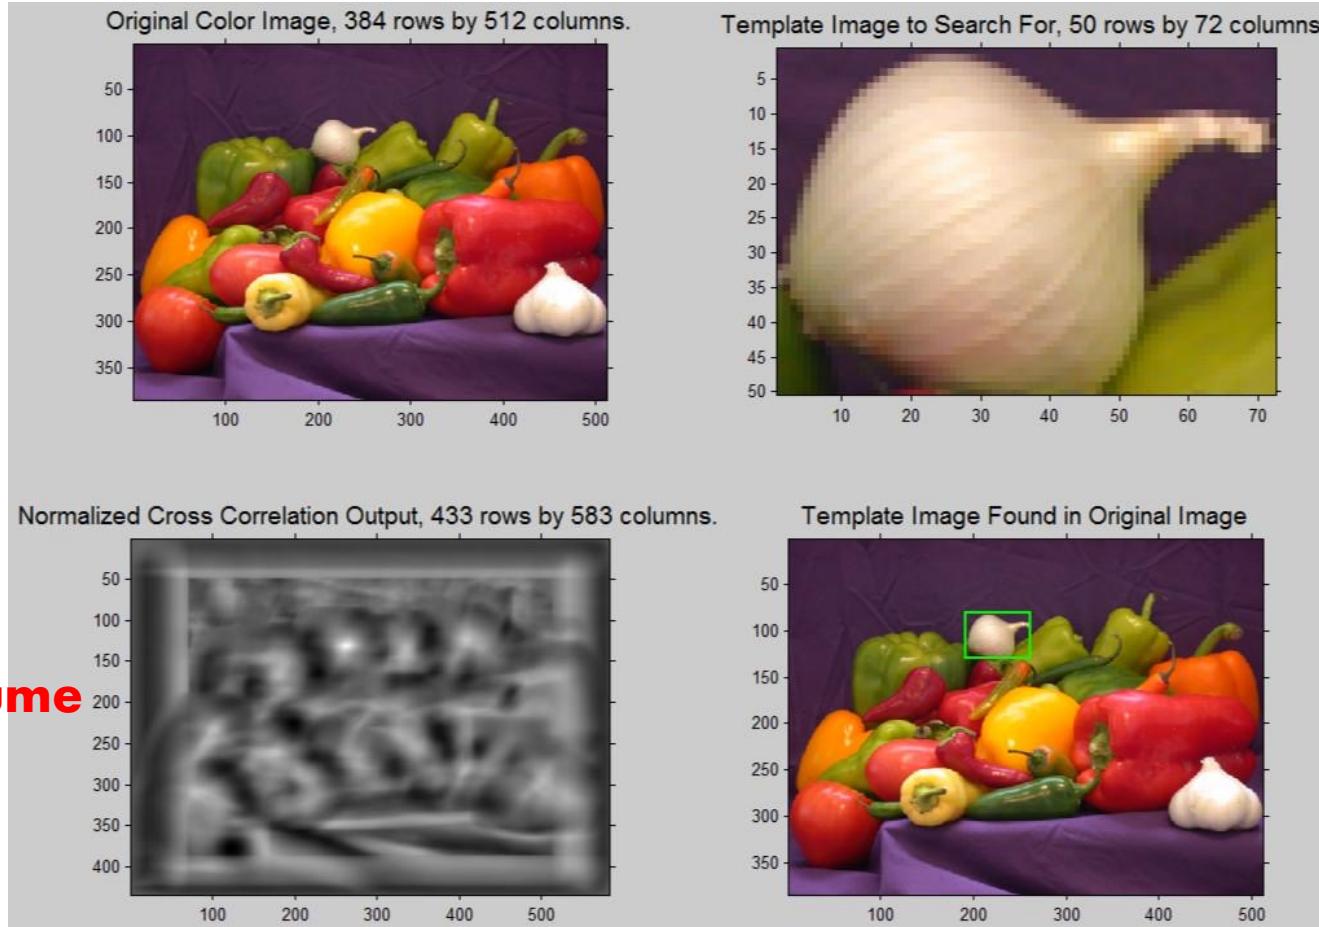
Supervised Optical Flow using traditional principles

Pyramidal Processing, Image Warping & Cost volume processing



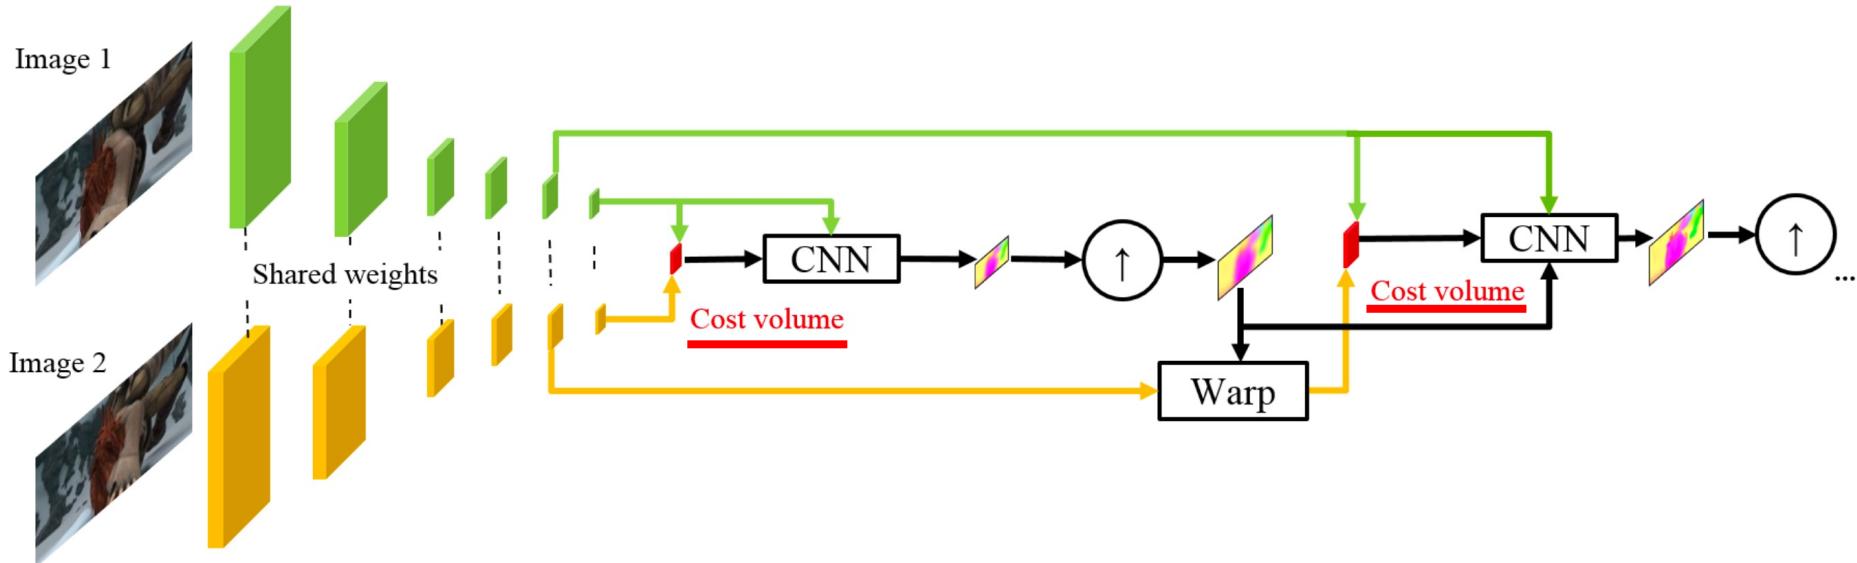
Deqing Sun, Xiaodong Yang, Ming-Yu Liu, and Jan Kautz. "PWC-Net: CNNs for Optical Flow Using Pyramid, Warping, and Cost Volume." CVPR 2019

Cost Volume



Supervised Optical Flow using traditional principles

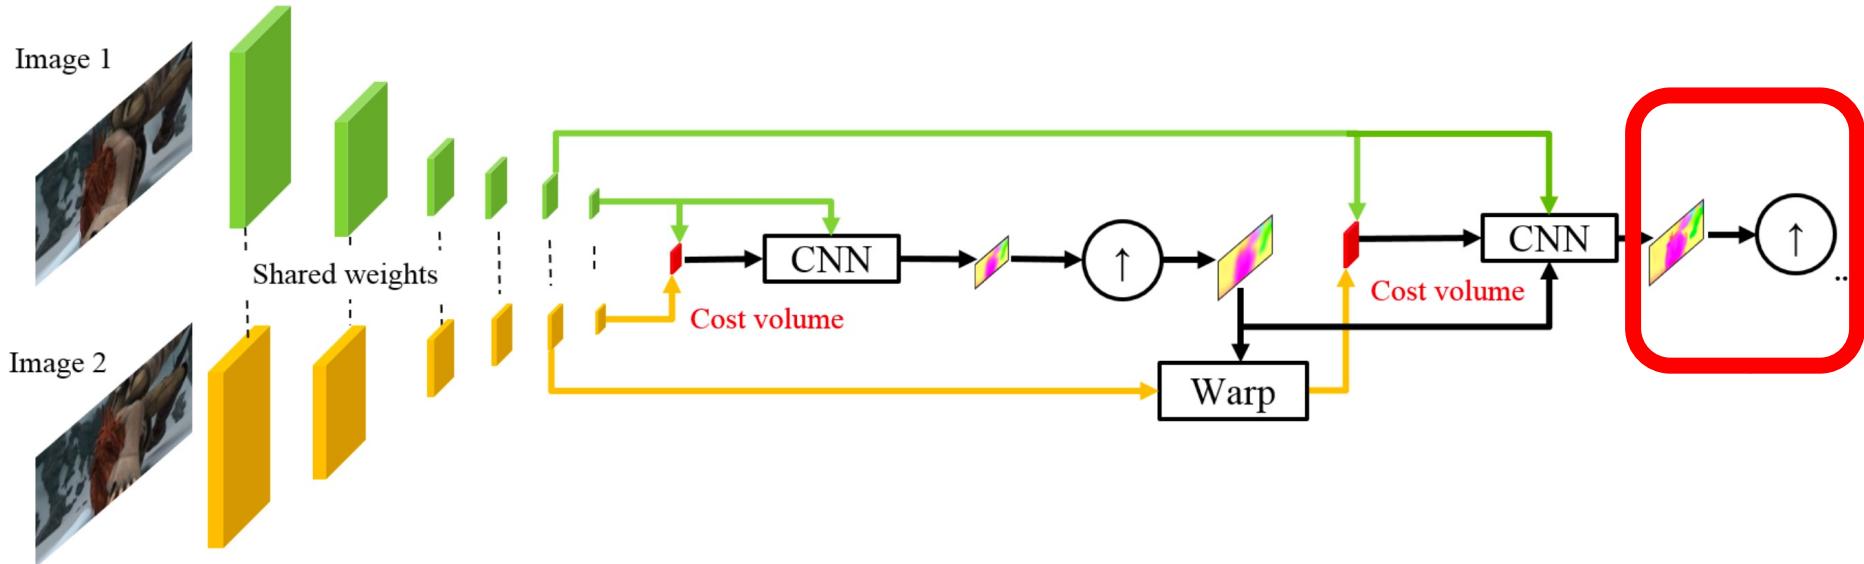
Pyramidal Processing, Image Warping & Cost volume processing



Deqing Sun, Xiaodong Yang, Ming-Yu Liu, and Jan Kautz. "PWC-Net: CNNs for Optical Flow Using Pyramid, Warping, and Cost Volume." CVPR 2018

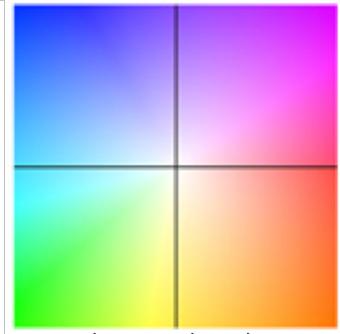
Supervised Optical Flow using traditional principles

Pyramidal Processing, Image Warping & Cost volume processing

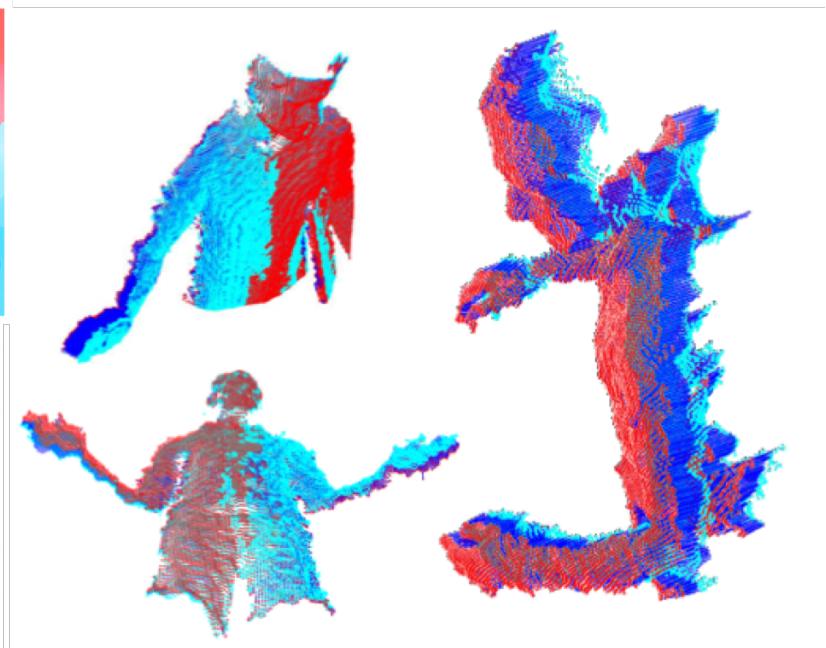


Deqing Sun, Xiaodong Yang, Ming-Yu Liu, and Jan Kautz. "PWC-Net: CNNs for Optical Flow Using Pyramid, Warping, and Cost Volume." CVPR 2018

Scene Flow

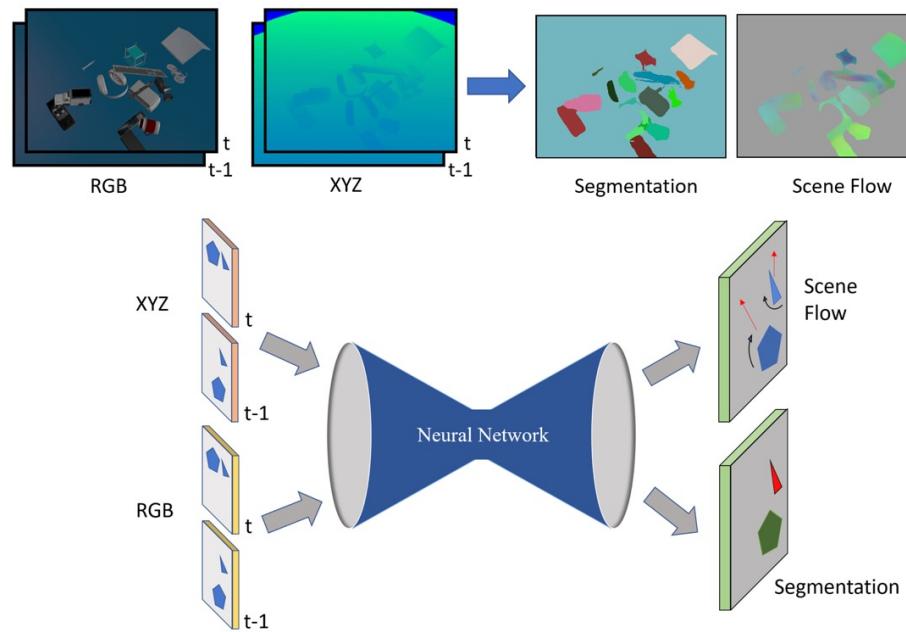


A Database and Evaluation Methodology for Optical Flow.
Baker et al. IJCV. 2011



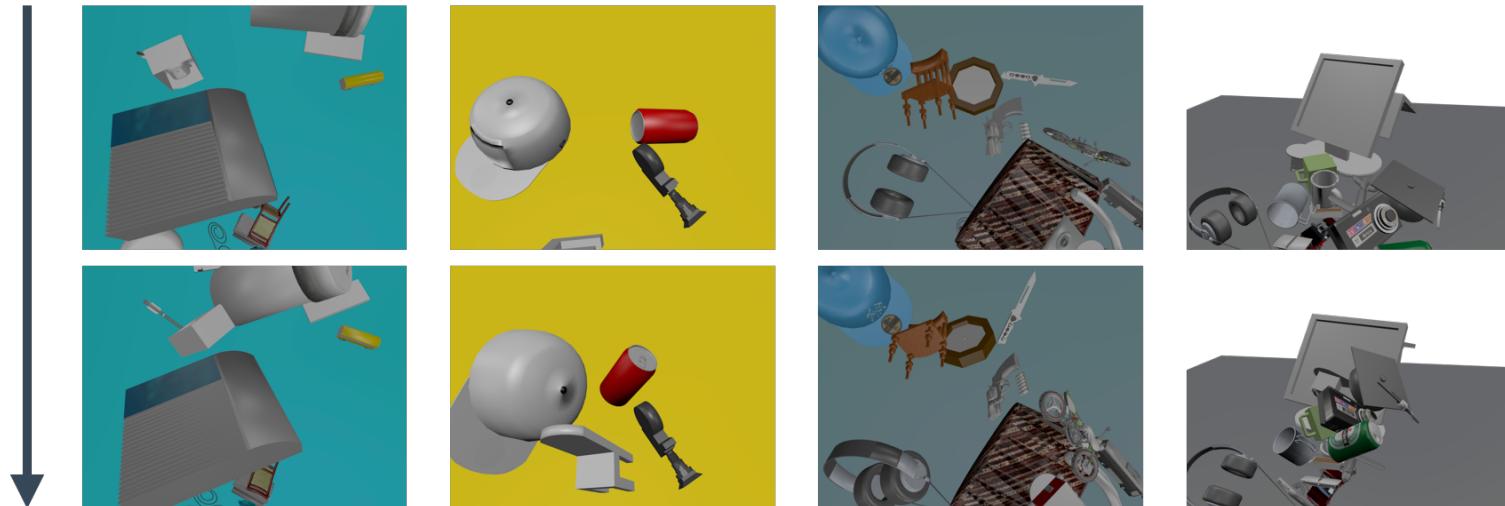
A Primal-Dual Framework for Real-Time Dense RGB-D Scene Flow. Jaimez et al. ICRA, 2015.

Motion-based Object Segmentation based on Dense RGB-D Scene Flow



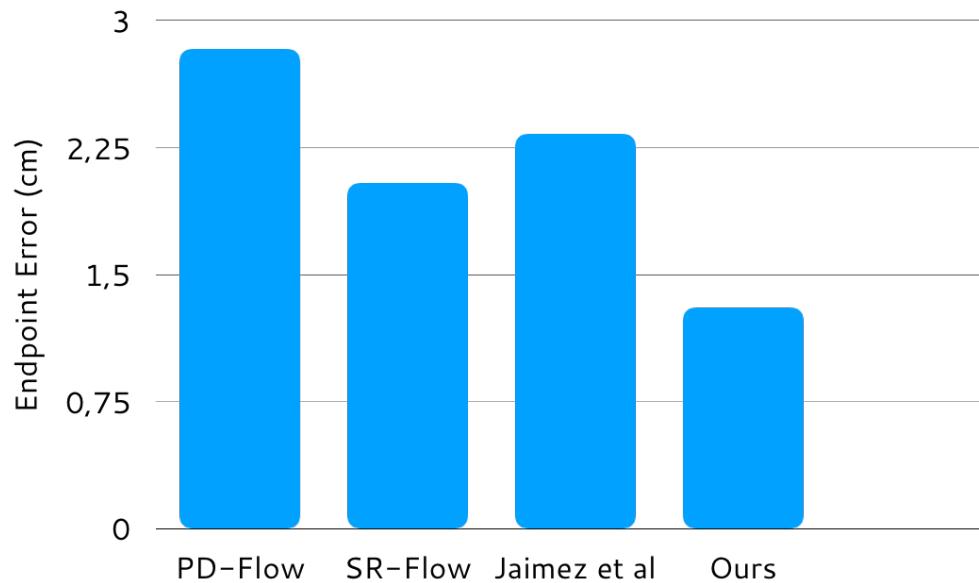
Motion-based Object Segmentation based on Dense RGB-D Scene Flow

Time

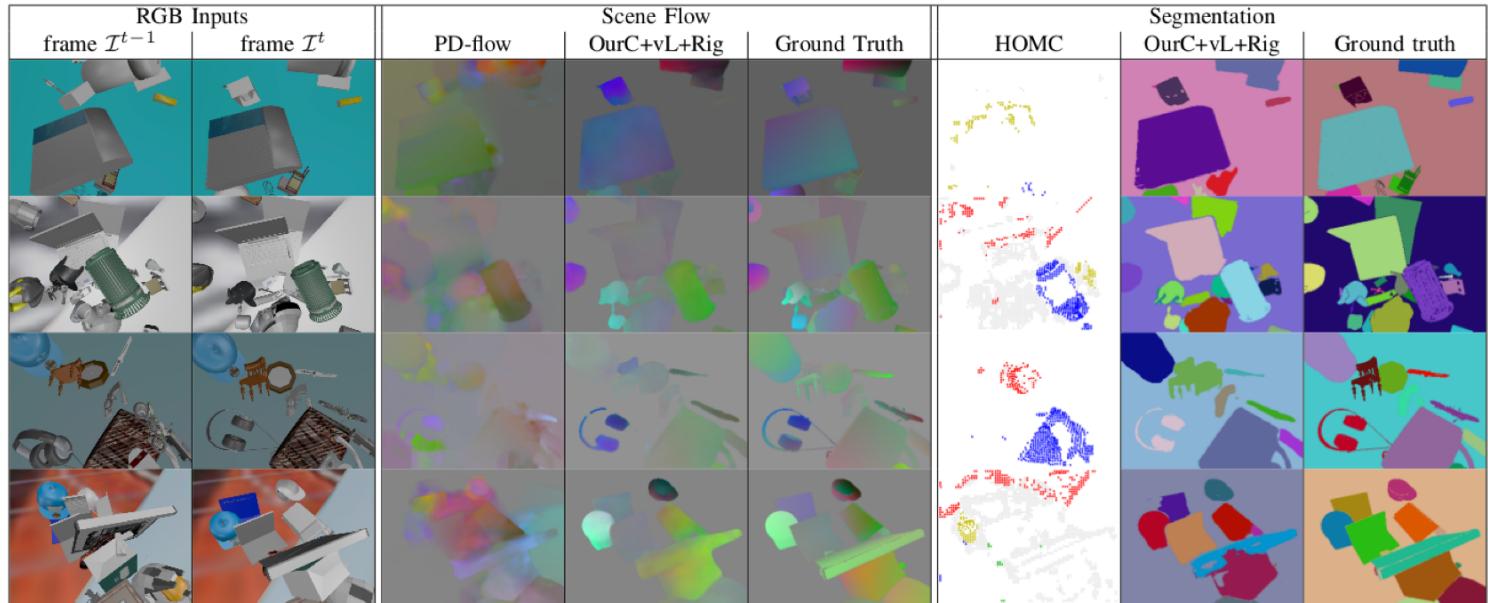


Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

Motion-based Object Segmentation based on Dense RGB-D Scene Flow

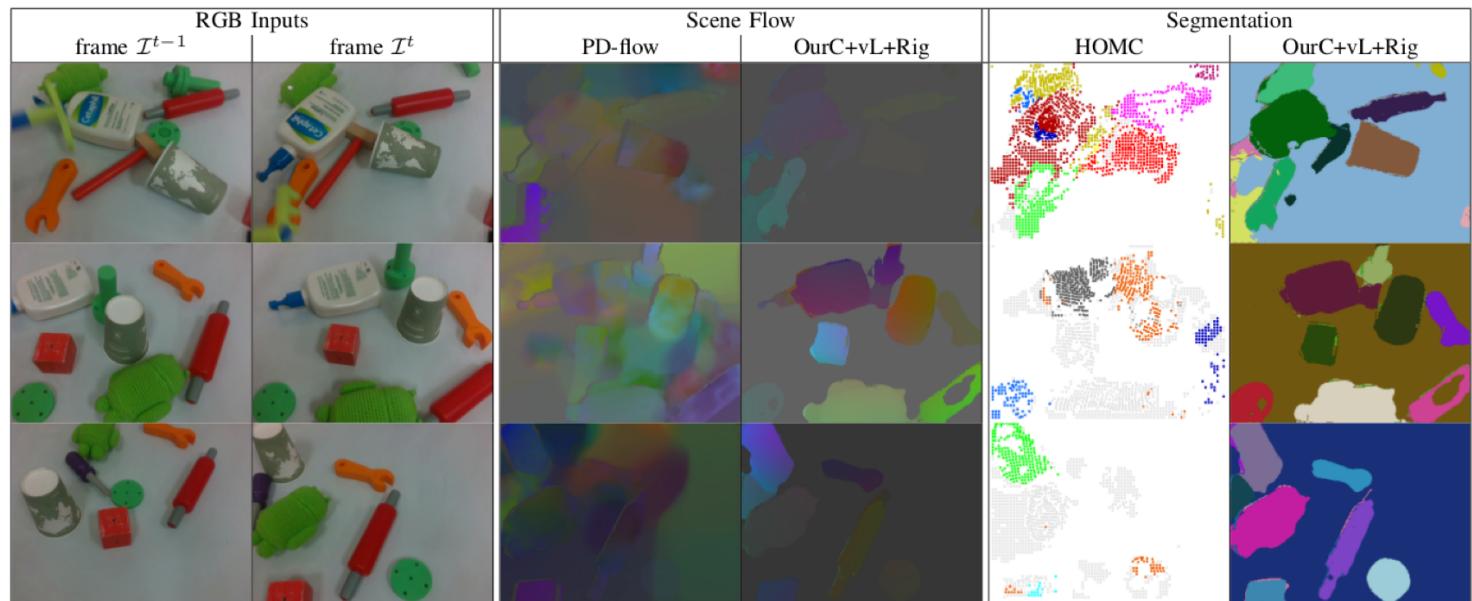


Motion-based Object Segmentation based on Dense RGB-D Scene Flow



Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

Motion-based Object Segmentation based on Dense RGB-D Scene Flow



Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

CS231

Introduction to Computer Vision



Next lecture:

Optimal Recursive Estimation