

CS231A

Computer Vision: From 3D Reconstruction to Recognition



Neural Radiance Fields II – Applications

Outline

- Recap: Neural Radiance Fields
- Applications
 - Pose Estimation
 - Motion-planning

The problem of novel view synthesis



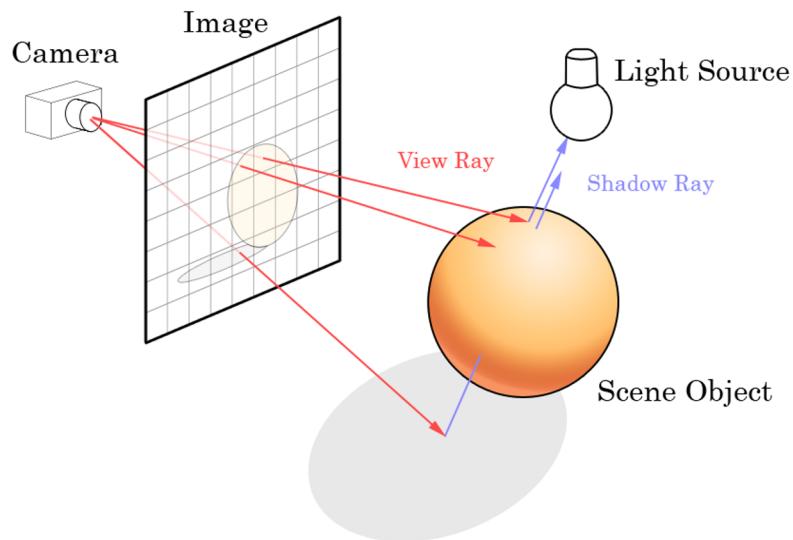
Inputs: sparsely sampled images of scene

Outputs: new views of same scene
(rendered by our method)

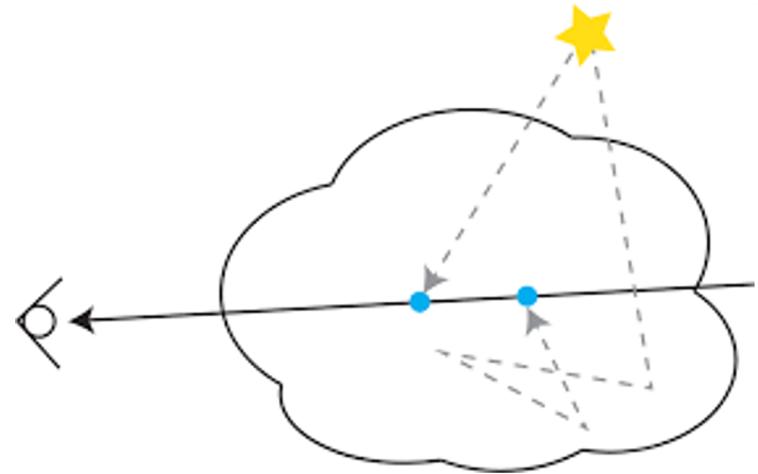
2

Mildenhall et al. ECCV 2020. <https://www.matthewtancik.com/nerf>

Ways to Render



Surface rendering



Volume rendering

Volume rendering equation

$$\mathbf{I}(D) = \mathbf{I}_0 T(0) + \int_0^D \mathbf{c}(s) \rho(s) T(s) ds$$

pixel color **radiance density**

$$T(s) = \exp \left(- \int_s^D \rho(t) dt \right)$$

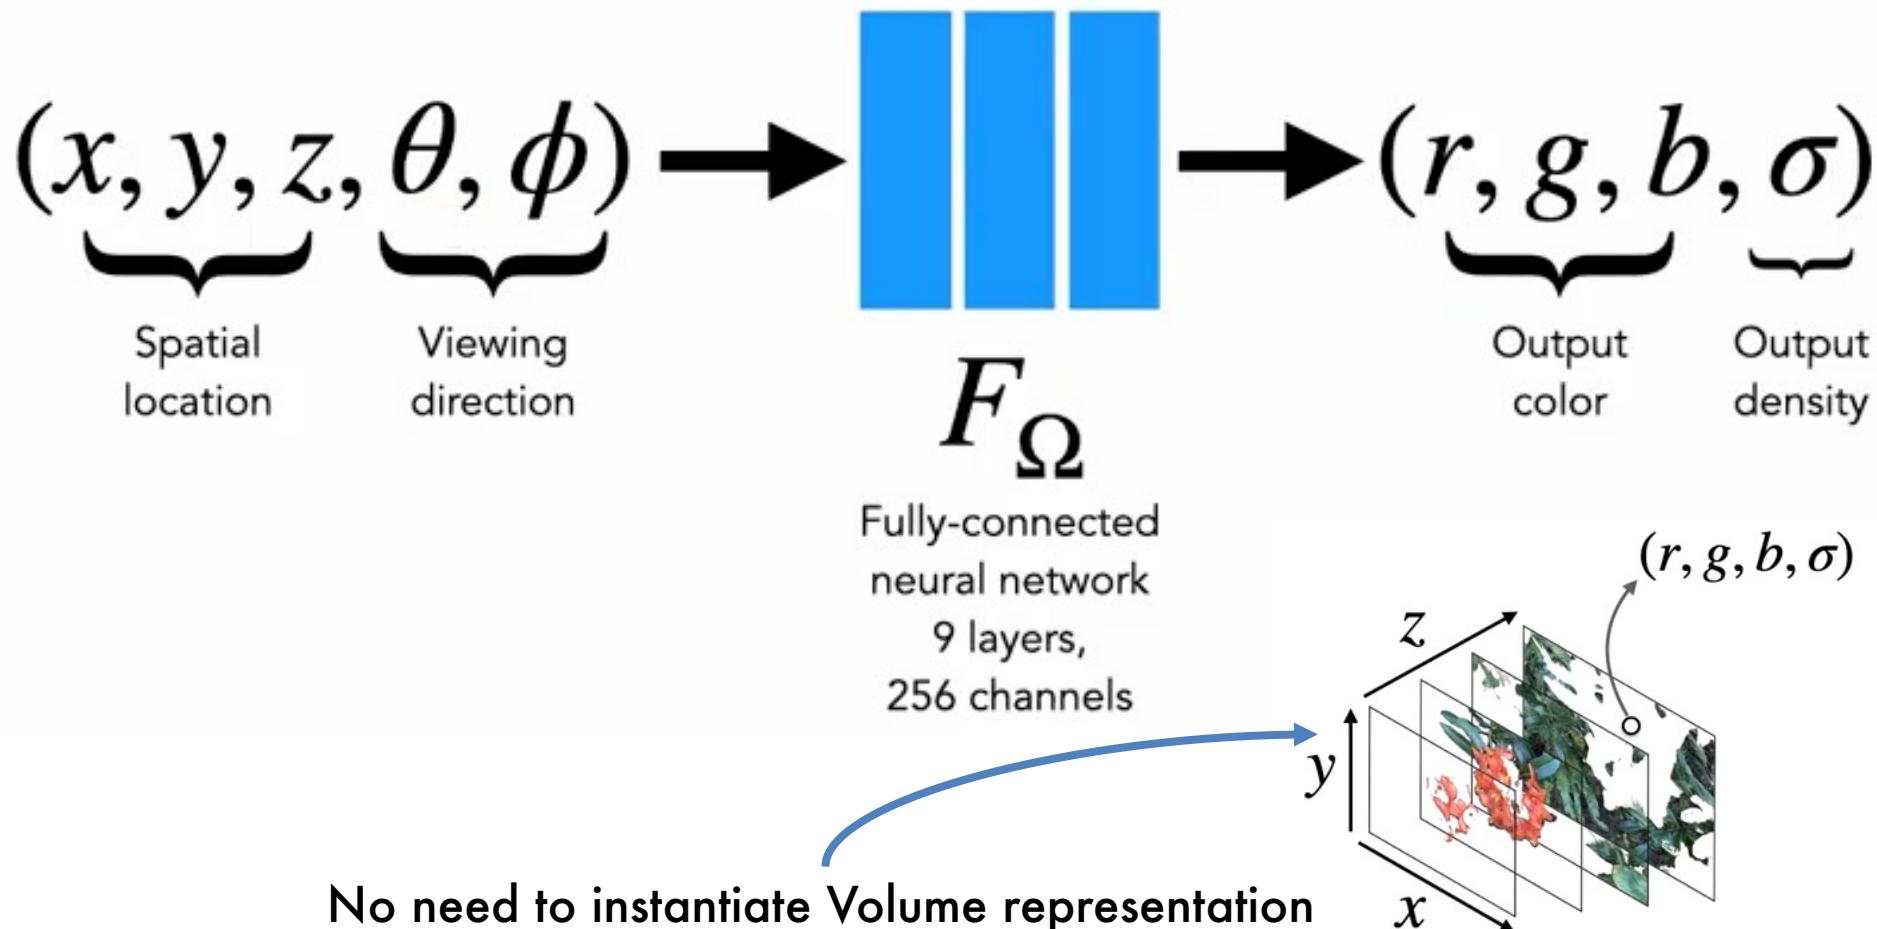
transparency

$(\mathbf{c}(s), \rho(s))$

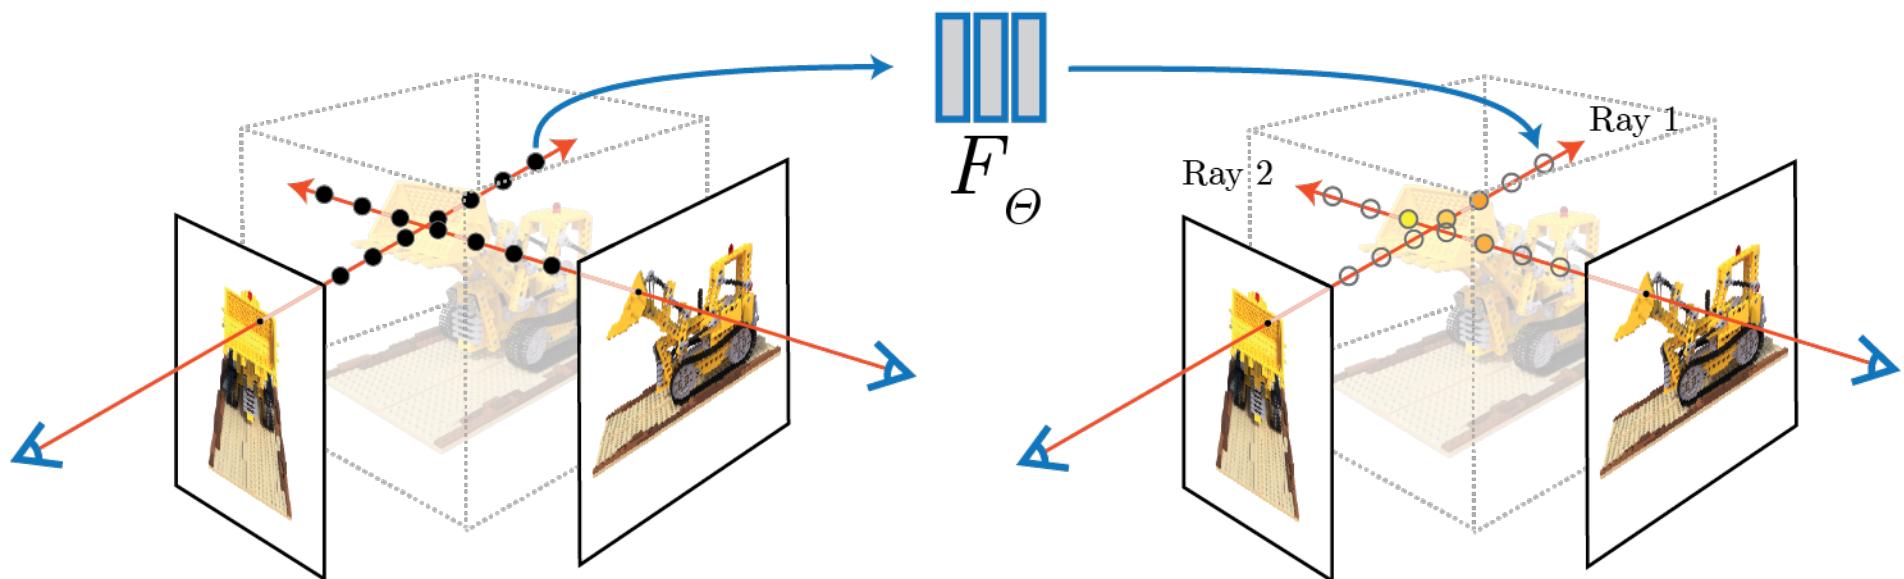
NeRF (neural radiance fields)

- Neural network as a volume representation using volume rendering to do view synthesis
- $(x, y, z, \theta, \phi) \rightarrow \text{color}, \text{opacity}$

Represent a scene as a continuous 5D function



Generate views with traditional volume rendering



Mildenhall et al. ECCV 2020. <https://www.matthewtancik.com/nerf>

Generate views with traditional volume rendering

Rendering model for ray $r(t) = o + td$:

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

weights colors

t = point along ray
C = Color of Pixel
c = color of point

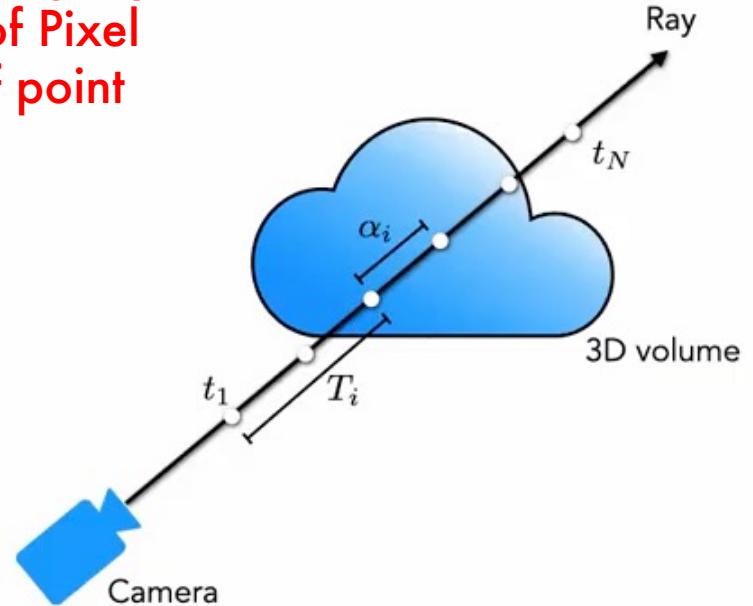
How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

Transparency

How much light is contributed by ray segment i :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i}$$



Function of segment length δt_i and volume density σ

From Presentation by Matthew Tancik: Neural Radiance Fields for View Synthesis. 2020.

Volume rendering is trivially differentiable

Rendering model for ray $r(t) = o + td$:

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

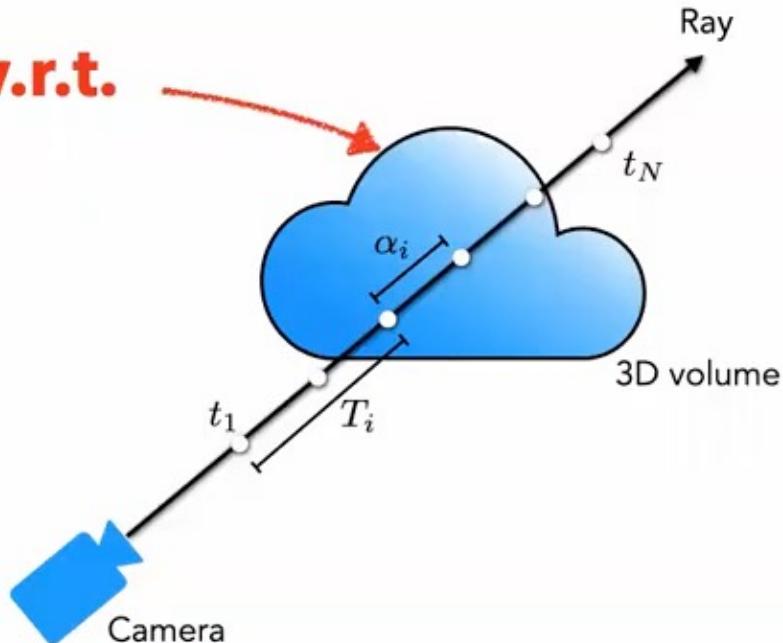
weights colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

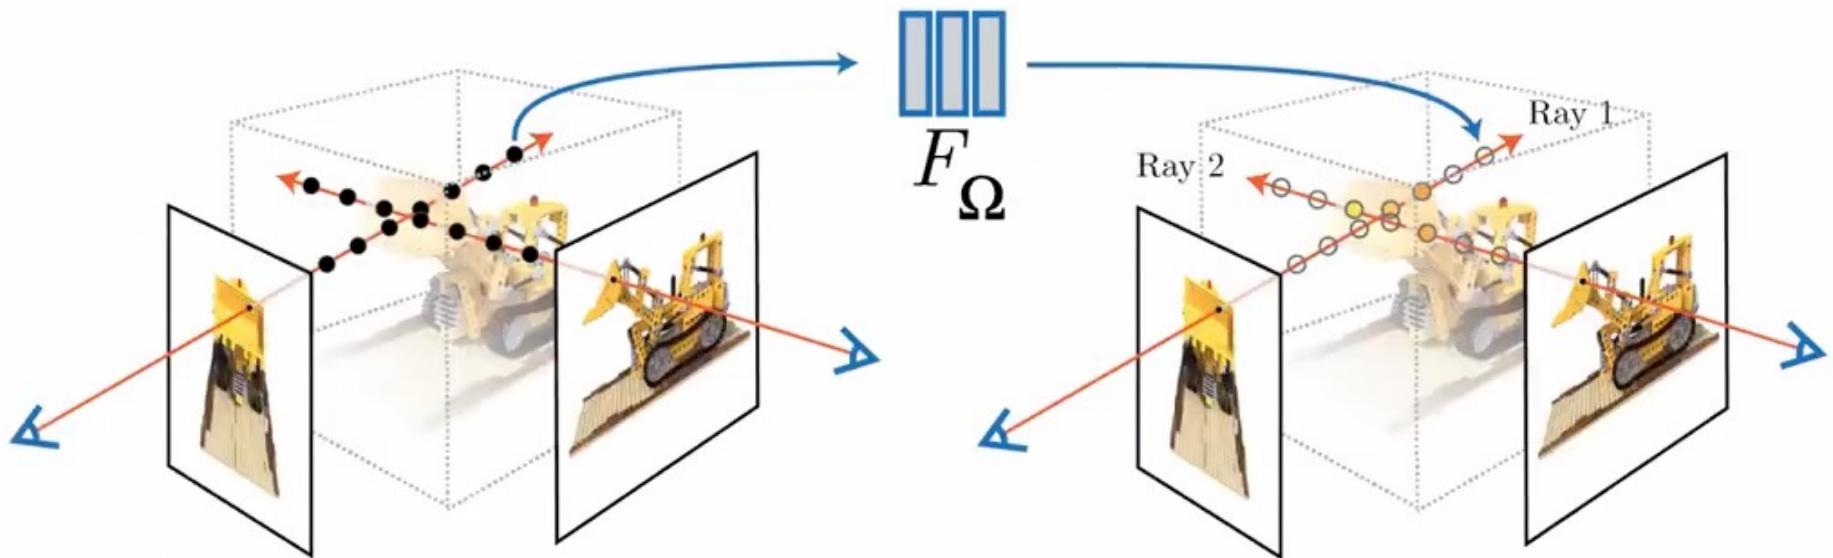
How much light is contributed by ray segment i :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i}$$

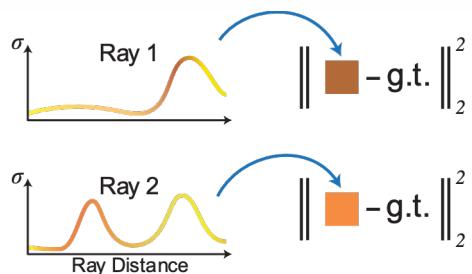


From Presentation by Matthew Tancik: Neural Radiance Fields for View Synthesis. 2020.

Optimize with gradient descent on rendering loss



$$\min_{\Omega} \sum_i \|\text{render}^{(i)}(F_{\Omega}) - I_{gt}^{(i)}\|^2$$



From Presentation by Matthew Tancik: Neural Radiance Fields for View Synthesis. 2020.

Training network to reproduce all input views of the scene



From Presentation by Matthew Tancik: Neural Radiance Fields for View Synthesis. 2020.

Results – Synthetic data



Results – Visualization Geometry



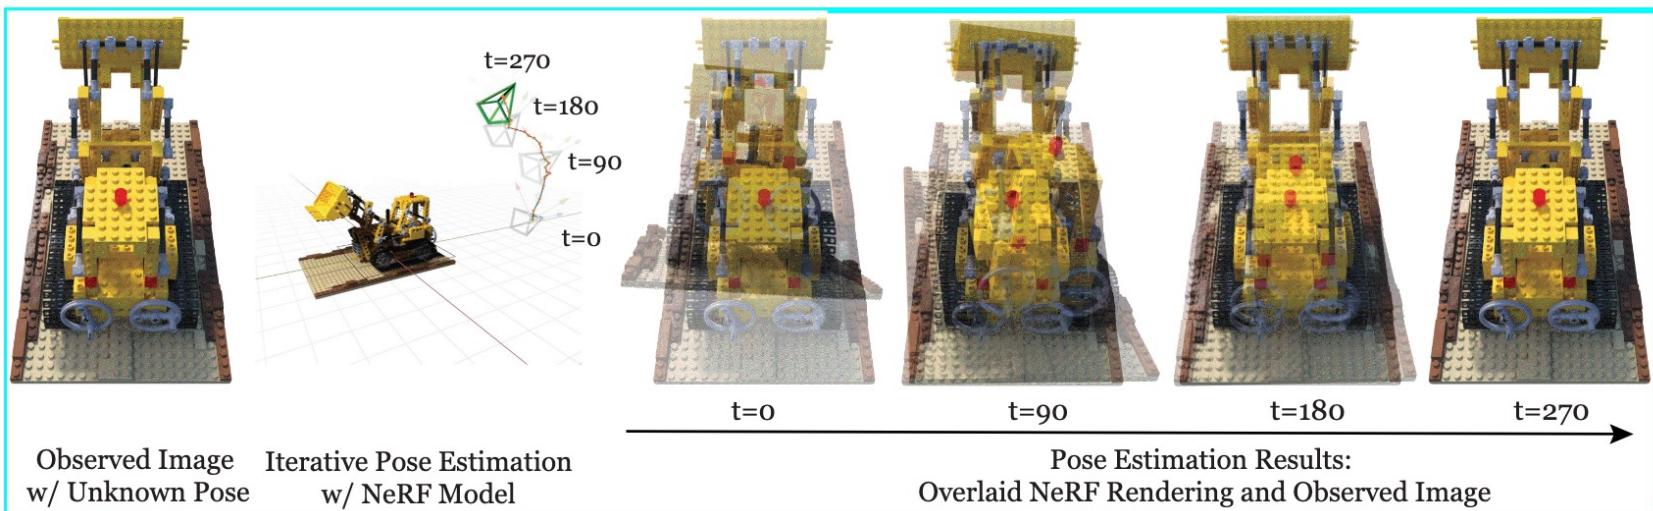
PollEverywhere

NeRF for Pose Estimation

iNeRF: Inverting Neural Radiance Fields for Pose Estimation

Lin Yen-Chen^{1,2}
Alberto Rodriguez²
¹Google Research

Pete Florence¹ Jonathan T. Barron¹
Phillip Isola² Tsung-Yi Lin¹
²Massachusetts Institute of Technology

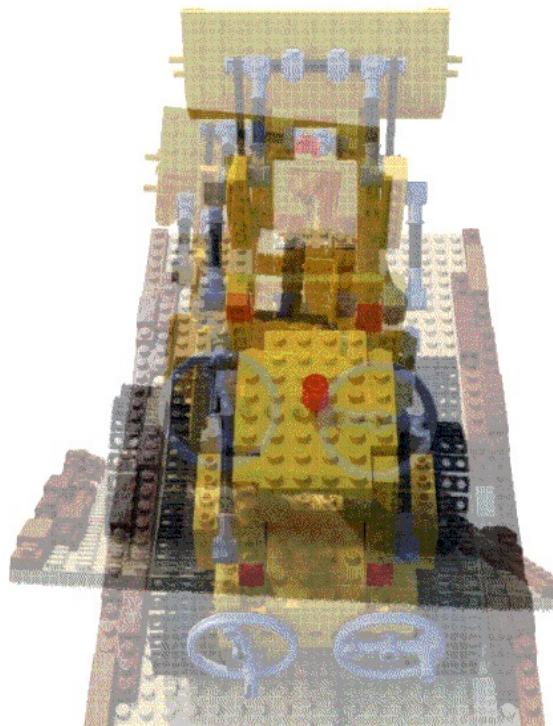
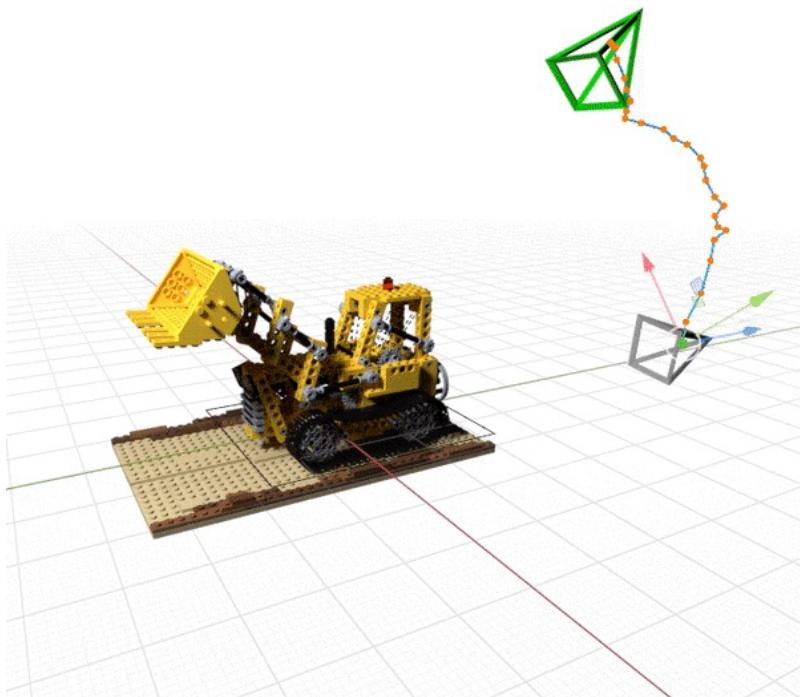


Input

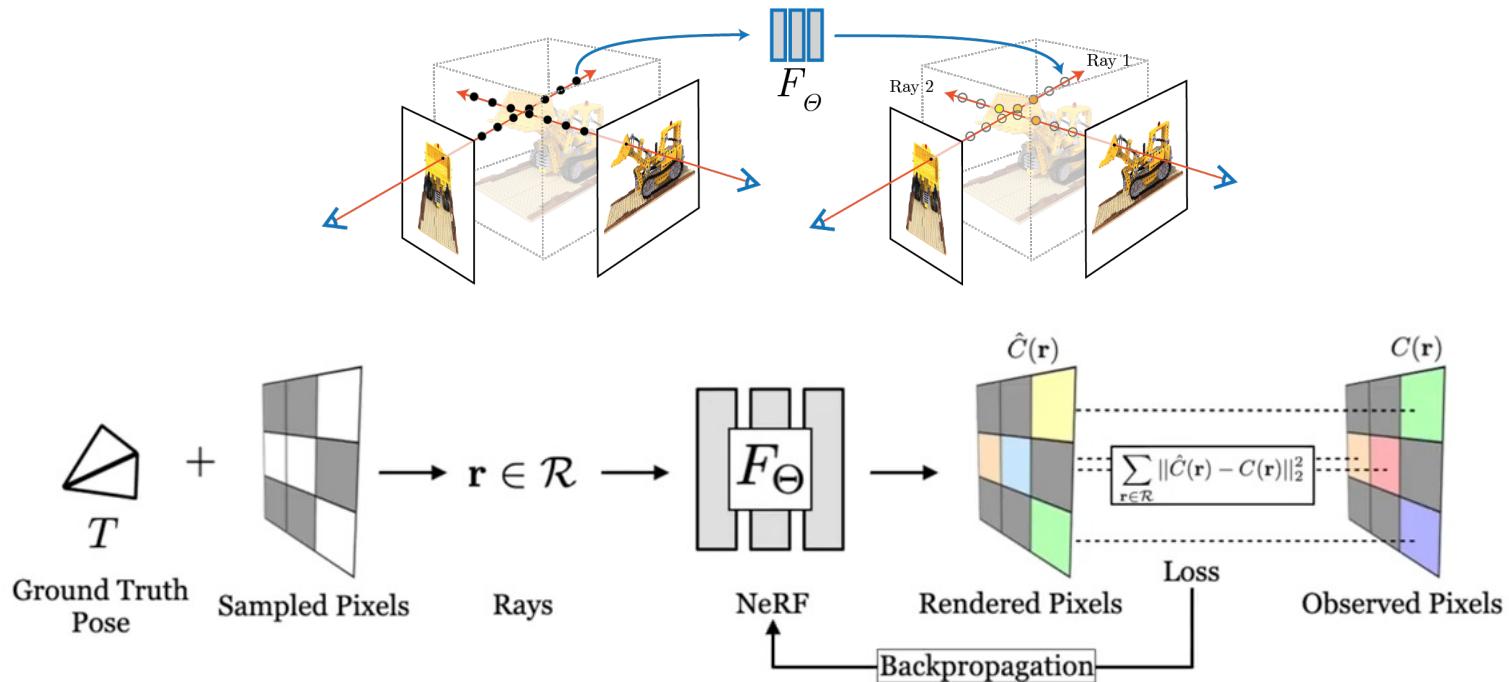
- An RGB Image
- Initial Pose
- NeRF model
- Not required
 - Depth
 - Mesh model

Output

Position and orientation of Camera relative to object. (6DoF)



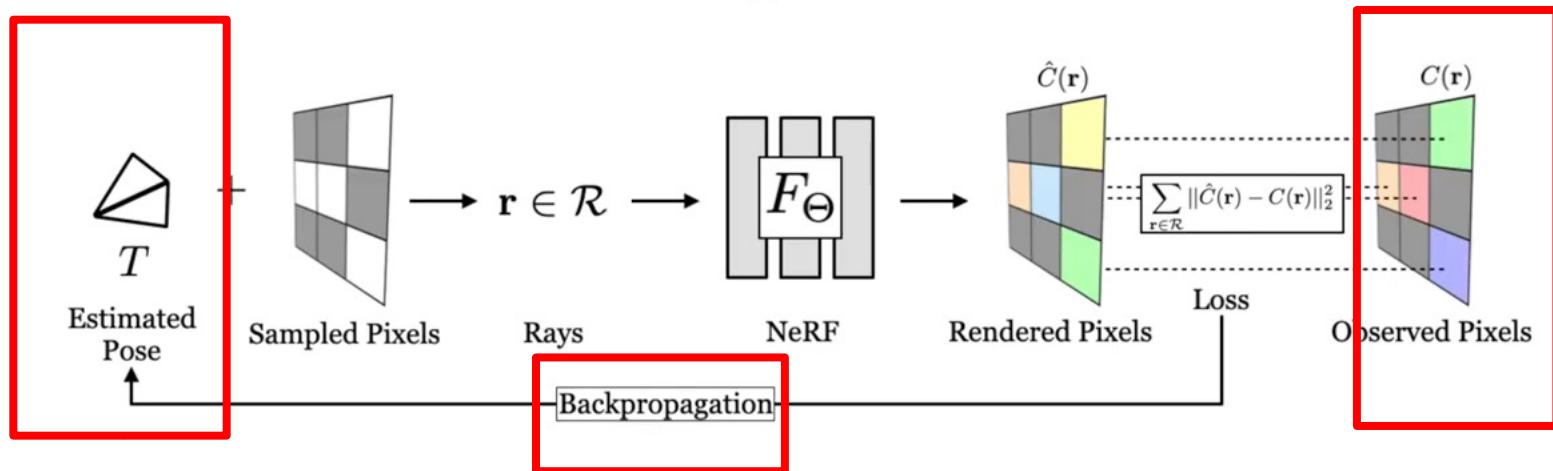
Recall: Training Vanilla NeRF



Synthesize Novel views

iNeRF: Inverting NeRF for Pose Estimation

$$\hat{T} = \underset{T \in \text{SE}(3)}{\operatorname{argmin}} \mathcal{L}(T \mid I, \Theta)$$



Gradient-based SE(3) Optimization

Loss at Current Optimization Iteration $L(\hat{T}_i | I, \theta)$

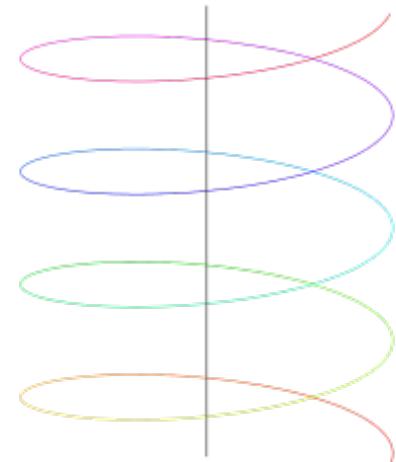
Parameterize \hat{T}_i using exponential coordinates to ensure solution lies on SE(3) manifold

With initial pose estimate \hat{T}_0 $\hat{T}_i = e^{[\mathcal{S}_i]\theta_i} \hat{T}_0,$

$$e^{[\mathcal{S}]\theta} = \begin{bmatrix} e^{[\omega]\theta} & K(\mathcal{S}, \theta) \\ 0 & 1 \end{bmatrix}$$

Screw axis $S = [\omega, v]$

Magnitude θ



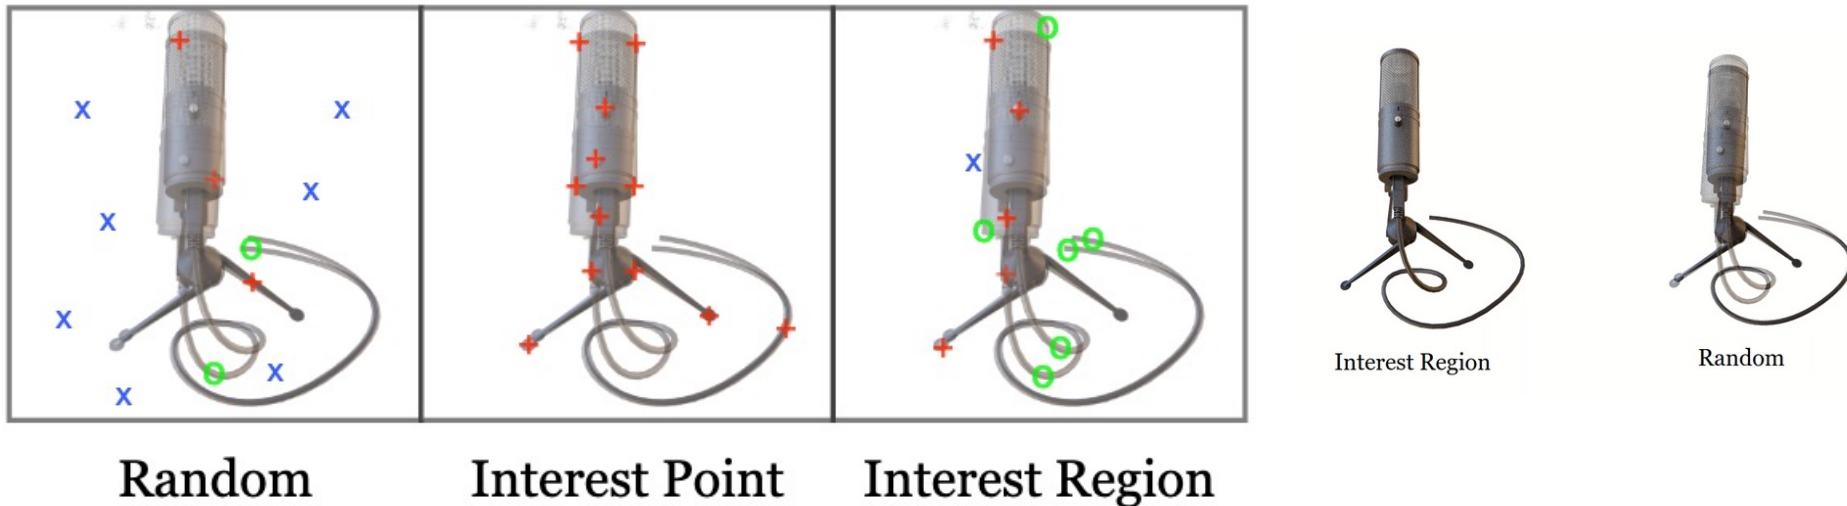
$$K(S, \theta) = (I\theta + (1 - \cos\theta)[\omega] + (\theta - \sin\theta)[\omega]^2)v$$

For full derivation see Modern Robotics by Lynch & Park, Section 3.3.3.1

Final Loss

$$\widehat{\mathcal{S}\theta} = \underset{S\theta \in \mathbb{R}^6}{\operatorname{argmin}} \mathcal{L}(e^{[\mathcal{S}]\theta} T_0 | I, \Theta)$$

Sampling Rays



✖ samples in background

✚ samples covered by rendered and observed images

○ samples covered only by either, the sampled or observed images

Interest Point Sampling - find key points.

Interest Region Sampling – Sampling from dilated masks around interest points

Results

Rotation: 45°

Category-Level Results



Input video.



Predicted pose & image.



Results on Real World images



NeRF versus iNeRF

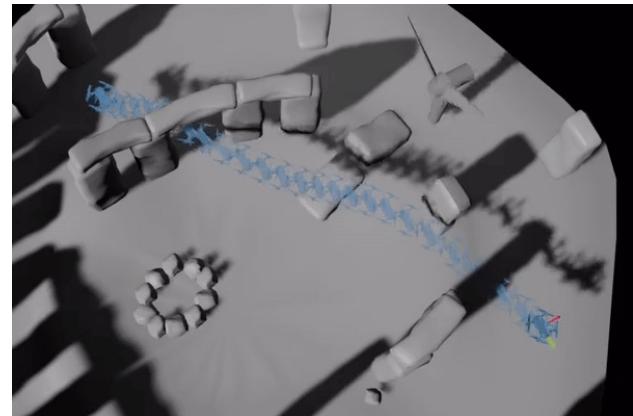
Method	NeRF	iNeRF
Input	Camera Pose	Image
Output	Image	Camera Pose

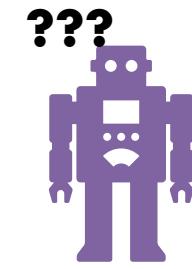
Outline

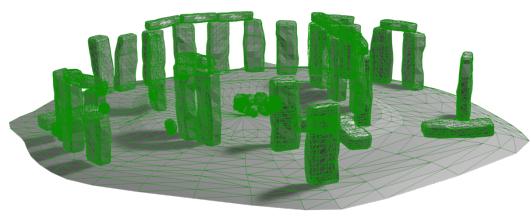
- Recap: Neural Radiance Fields
- Applications
 - Pose Estimation
 - Motion-planning

Vision-Only Robot Navigation in a Neural Radiance World

Michał Adamkiewicz*, Timothy Chen*, Adam Caccavale, Rachel Gardner,
Preston Culbertson, Jeannette Bohg, Mac Schwager
ICRA'22 + RAS Letters

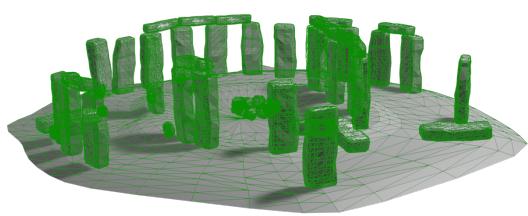




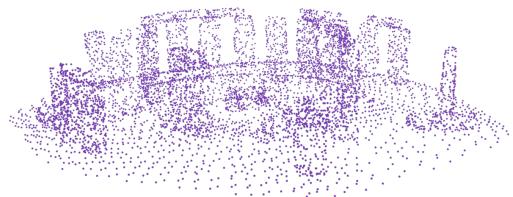


Mesh



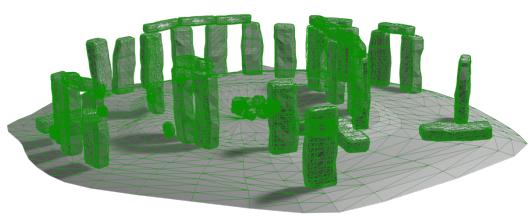


Mesh

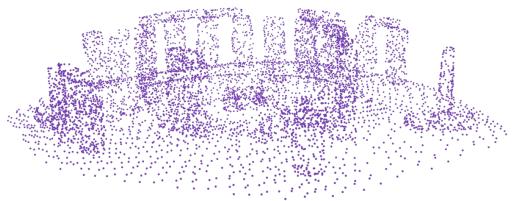


Point Cloud

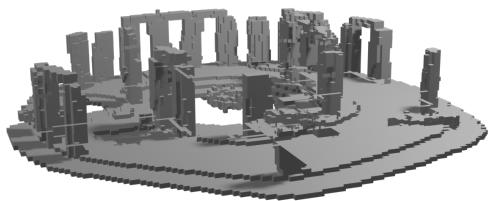




Mesh

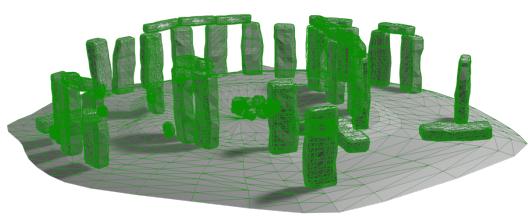


Point Cloud

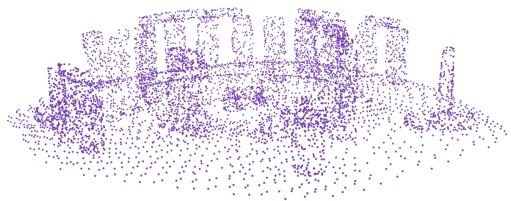


Voxels

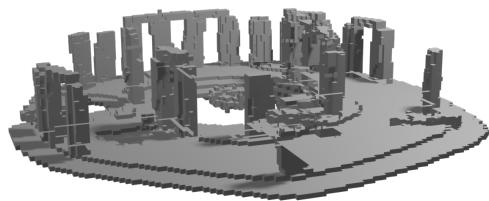




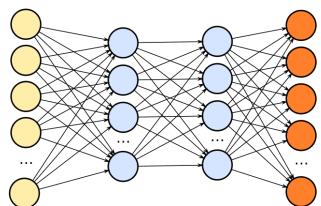
Mesh



Point Cloud

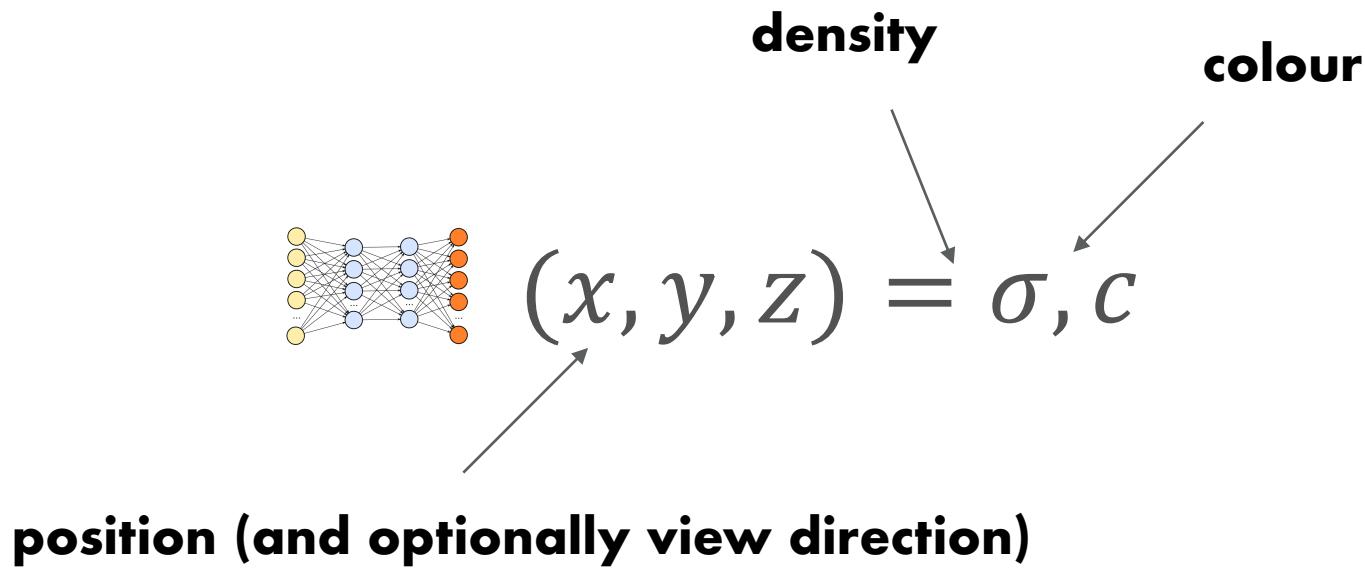


Voxels

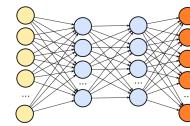


Implicit representations

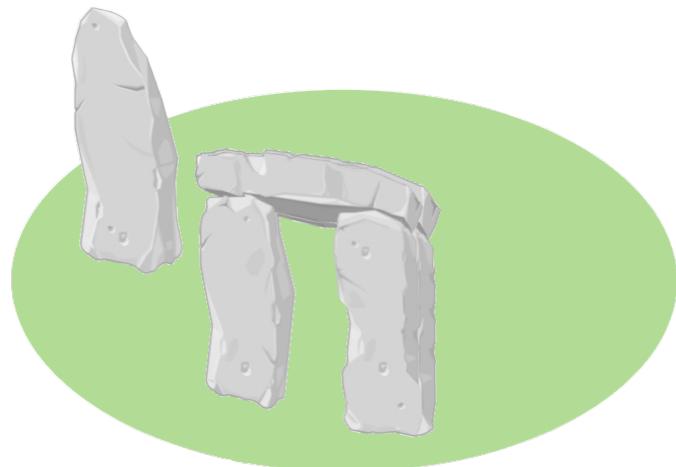
Implicit representations



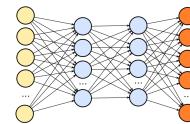
Implicit representations



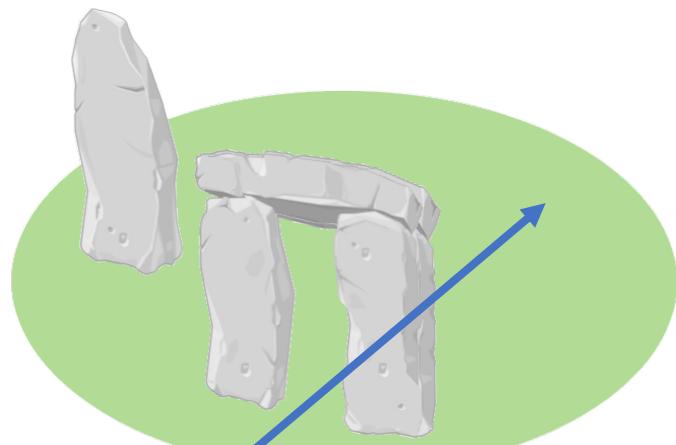
$$(x, y, z) = \sigma, c$$



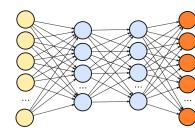
Implicit representations



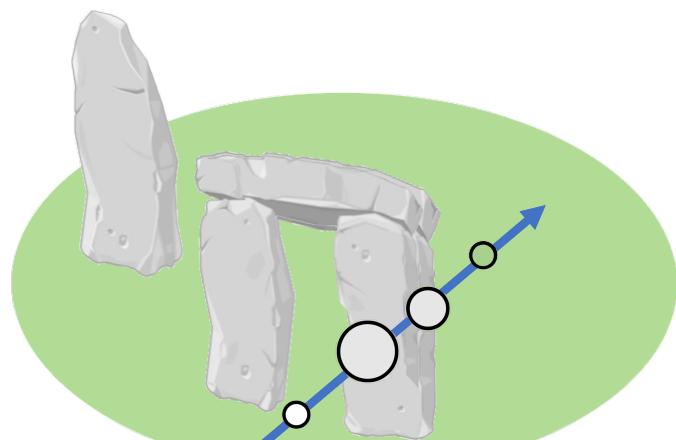
$$(x, y, z) = \sigma, c$$



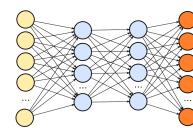
Implicit representations



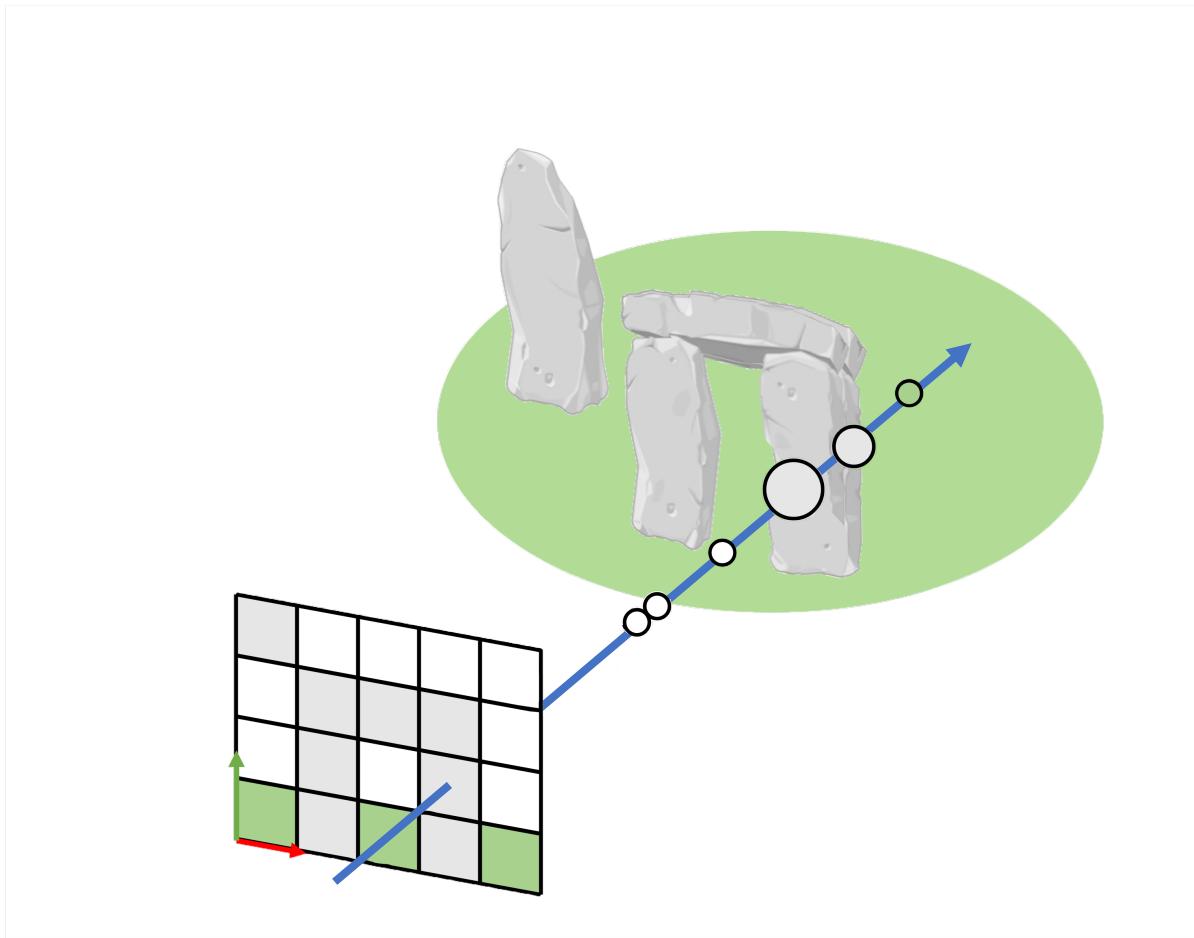
$$(x, y, z) = \sigma, c$$



Implicit representations



$$(x, y, z) = \sigma, c$$



$$\partial$$

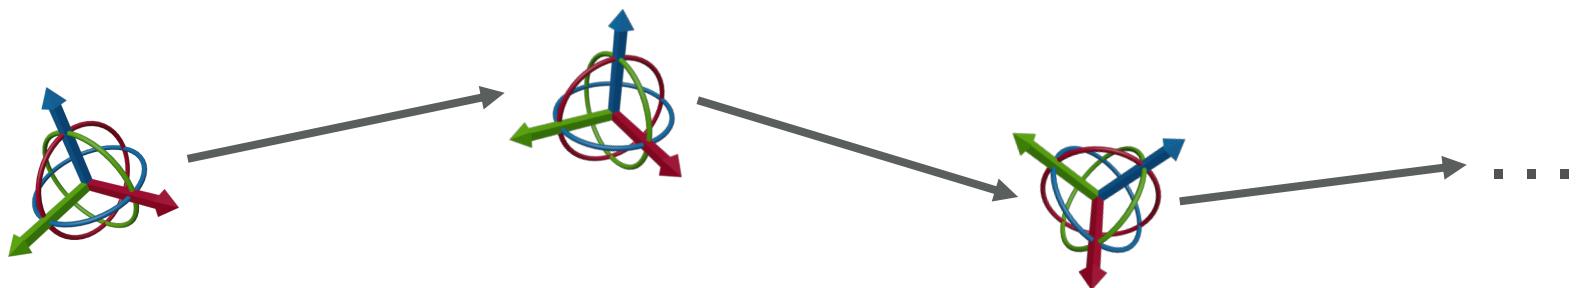
$$\bar{\partial}$$

$$\partial$$

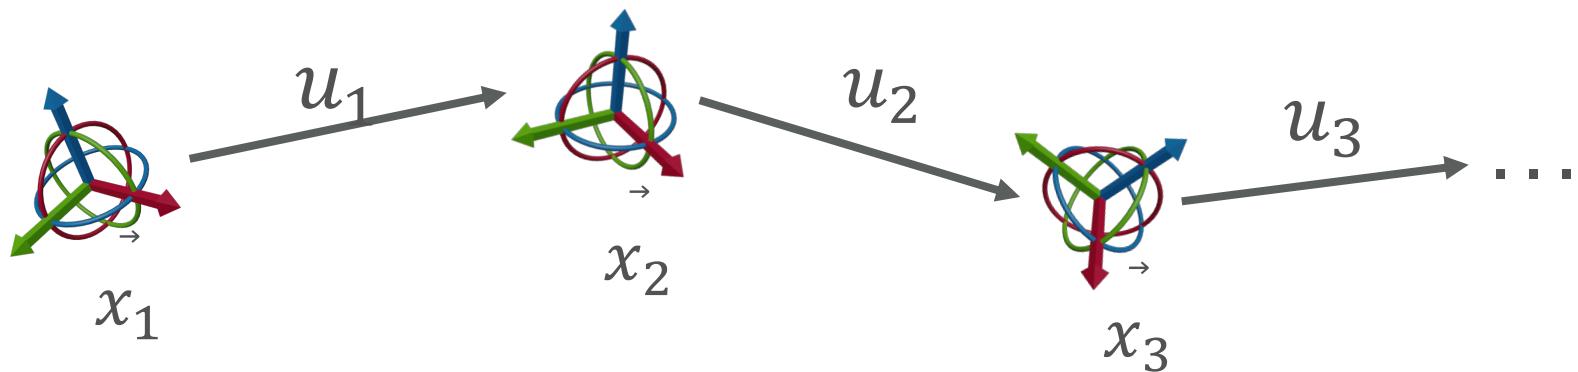
$$\bar{\partial}$$



Trajectory Planning

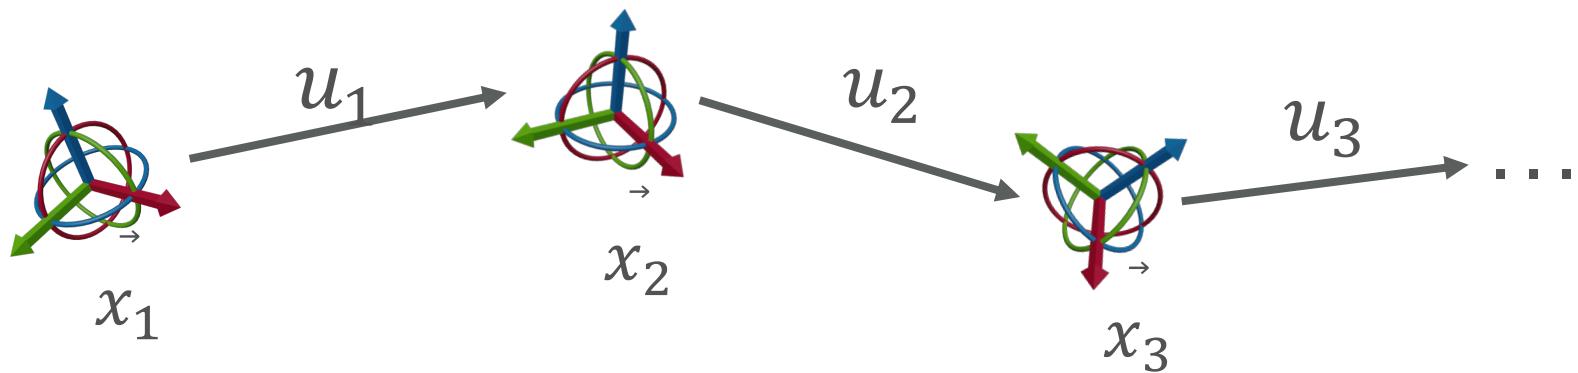


Trajectory Planning



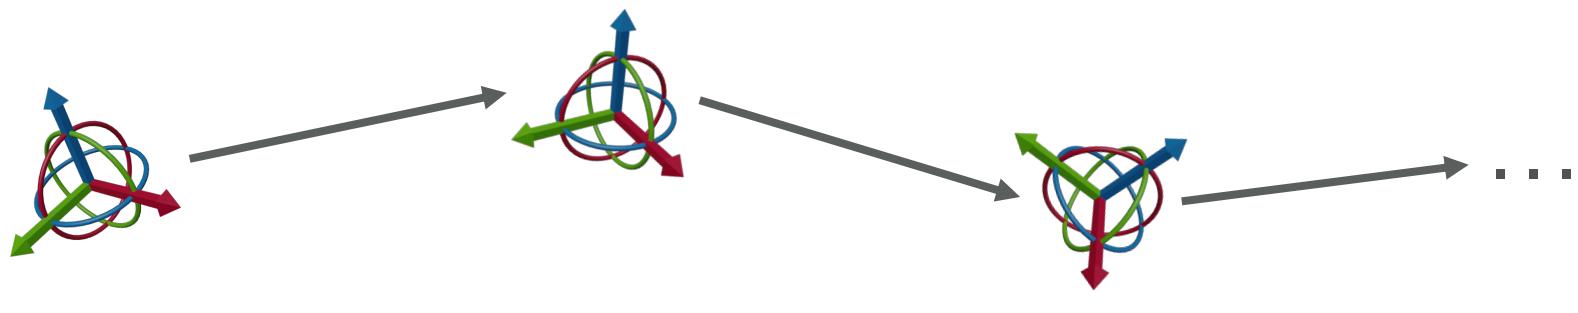
$$Loss = L_{collision} + L_{control effort}$$

Trajectory Planning



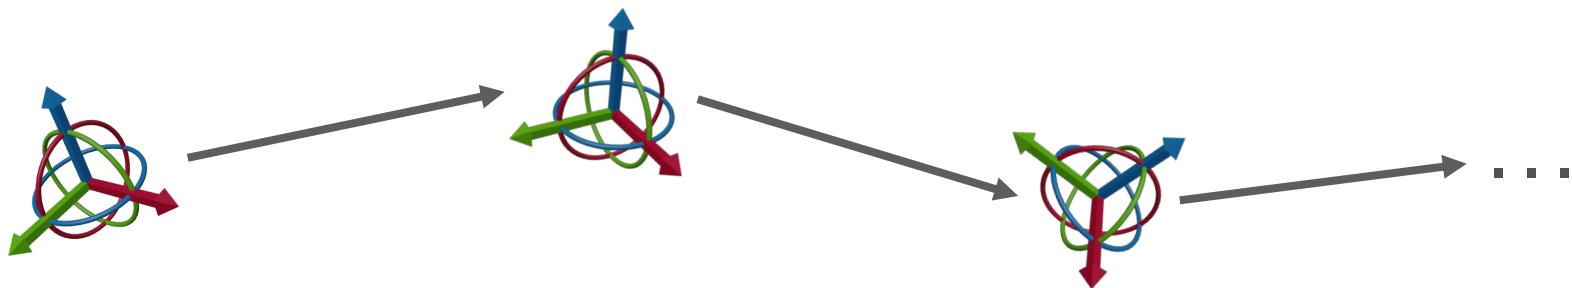
$$Loss = L_{collision} + L_{control effort}(u_1) + L_{control effort}(u_2) + \dots$$

Trajectory Planning



$$Loss = \text{Loss Function}(\text{Initial State}) + \text{Loss Function}(\text{Intermediate State}) + \text{Loss Function}(\text{Final State}) + \dots + \sum L_{control effort}$$

Trajectory Planning

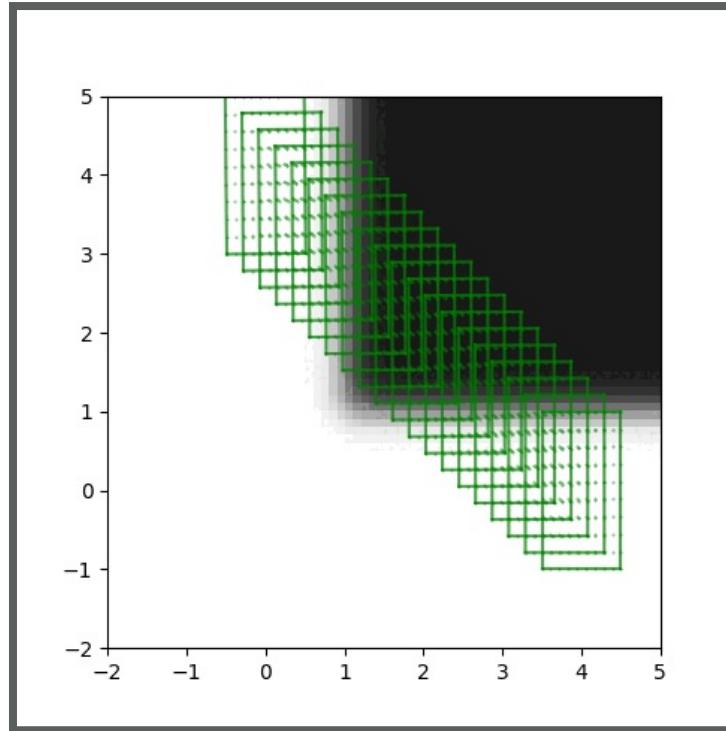


Loss =

$$|\vec{v}_1| \text{ ()} + |\vec{v}_2| \text{ ()} + |\vec{v}_3| \text{ ()} + \dots$$

$$+ \sum L_{control effort}$$

Trajectory Planning

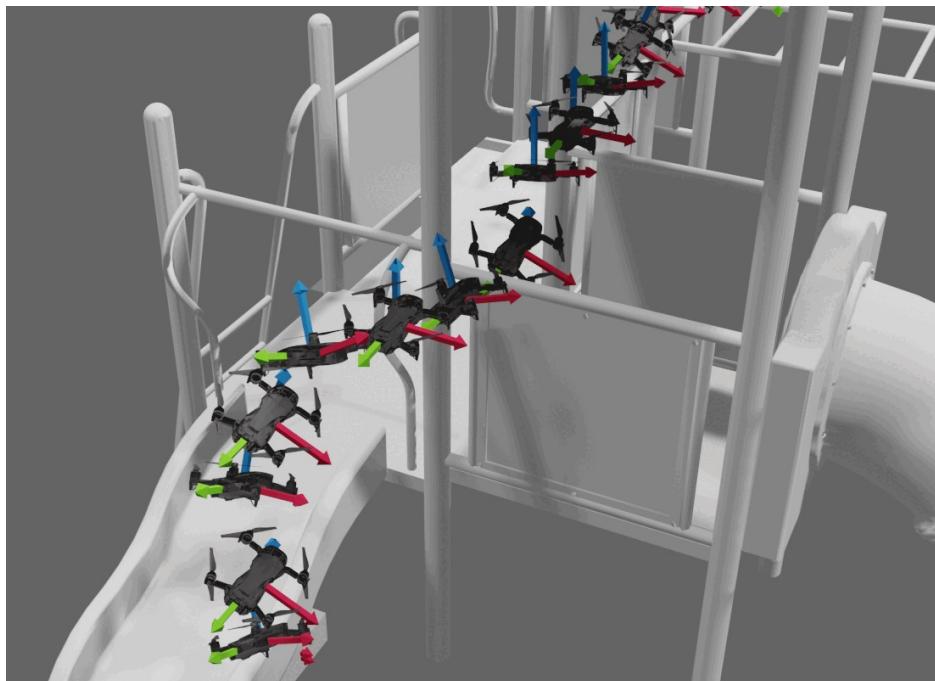


$Loss =$

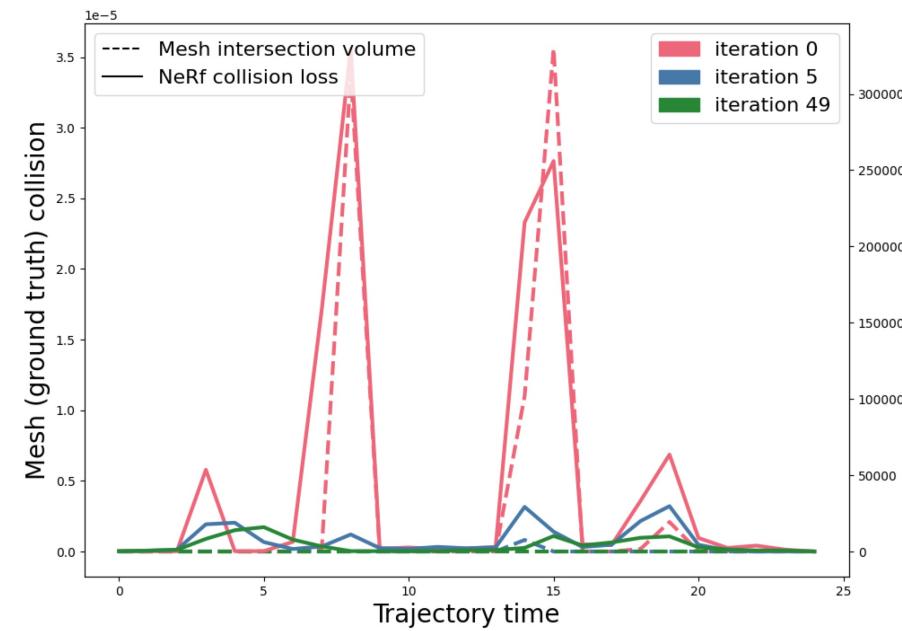
$$|\vec{v_1}| \text{ (Diagram of a neural network)} + |\vec{v_2}| \text{ (Diagram of a neural network)} + |\vec{v_3}| \text{ (Diagram of a neural network)} + \dots$$

$$+ \sum L_{control effort}$$

Trajectory Planning

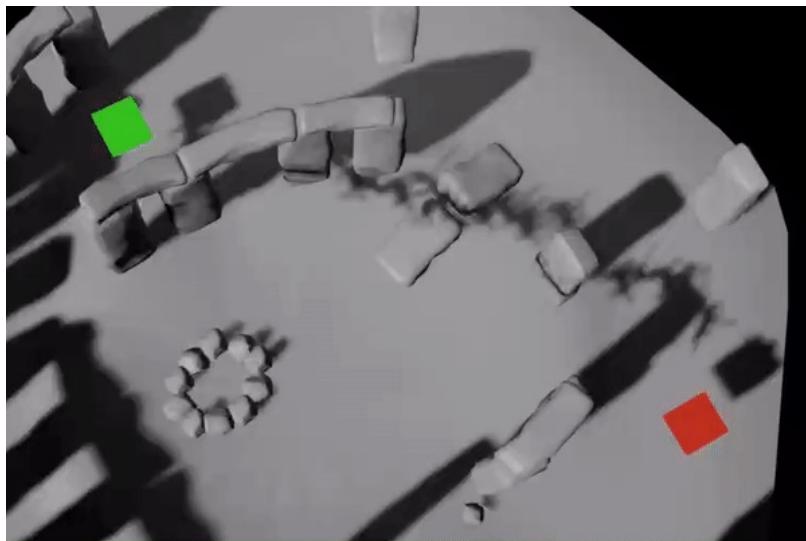


Optimisation process

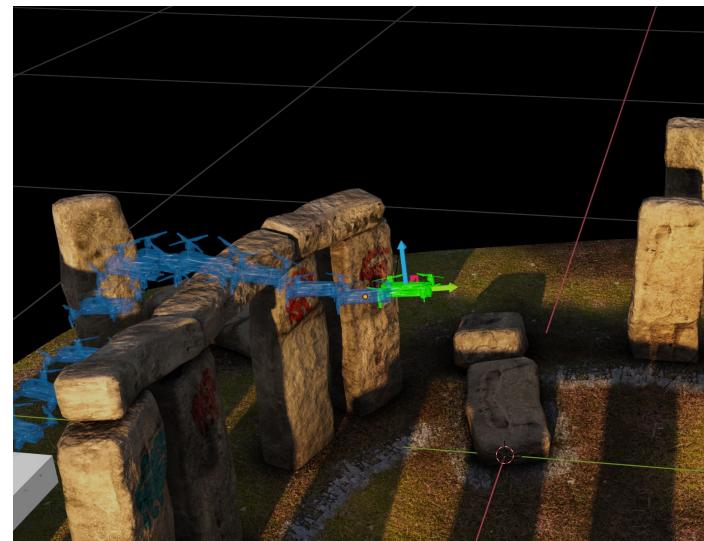


GT mesh comparison

Trajectory Planning



A^{*} initialisation



Avoids obstacles

Trajectory Planning



Aware of robot geometry



Aware of scene detail

State Estimation – Nonlinear System

$$\boldsymbol{\mu}_{t|t-1} = f(\boldsymbol{\mu}_{t-1}, \mathbf{u}_t)$$

$$\mathbf{A}_{t-1} = \frac{\partial f(\mathbf{x}, \mathbf{u}_t)}{\partial \mathbf{x}} \Big|_{\mathbf{x}=\boldsymbol{\mu}_{t-1}}$$

Predict Step

$$\boldsymbol{\Sigma}_{t|t-1} = \mathbf{A}_{t-1} \boldsymbol{\Sigma}_{t-1} \mathbf{A}_{t-1}^T + \mathbf{Q}_{t-1}$$

State Estimation

$$Loss = |\text{camera icon} - NeRF(\vec{x})|^2 + Loss_{Process}(\vec{x})$$



State Estimation

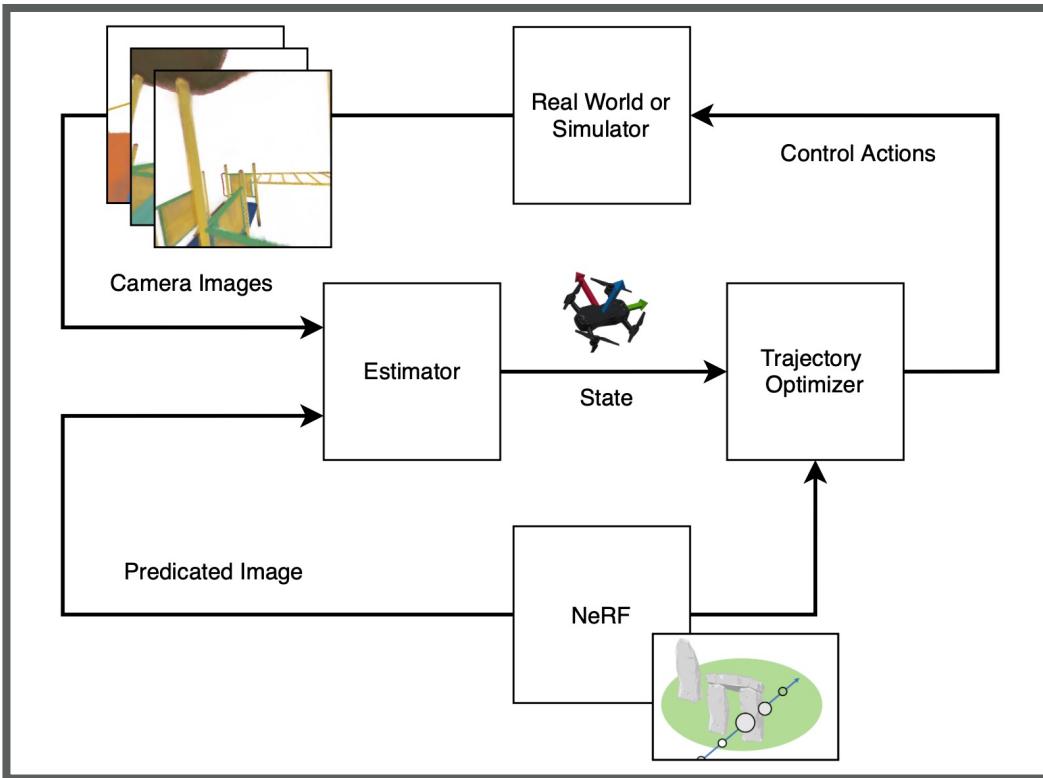
$$Loss = |\text{camera icon} - NeRF(\vec{x})|^2 + Loss_{Process}(\vec{x}, \Sigma_{t|t-1})$$

Minimizing this gives posterior state estimate

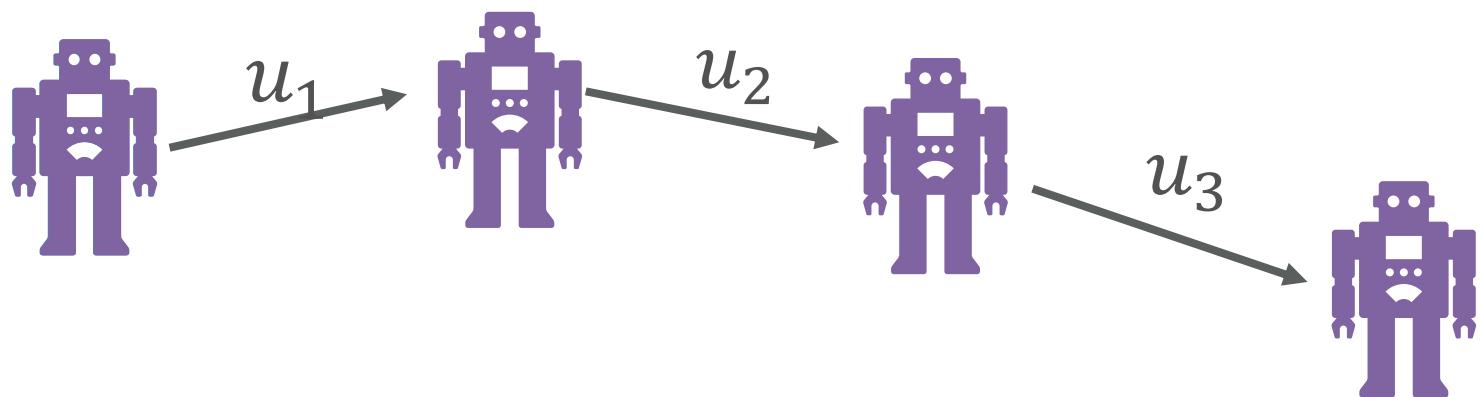
$$\Sigma_t = \left(\frac{\partial^2 Loss(x)}{\partial x^2} \Big|_{x=x_{opt}} \right)^{-1}$$



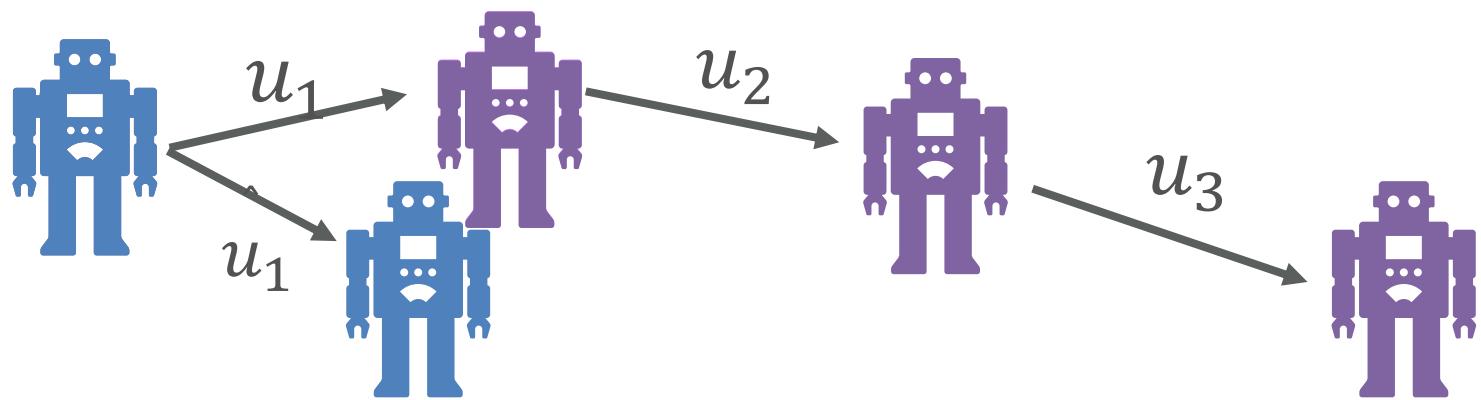
Full System Combining State Estimation and MPC controller



MPC controller



MPC controller

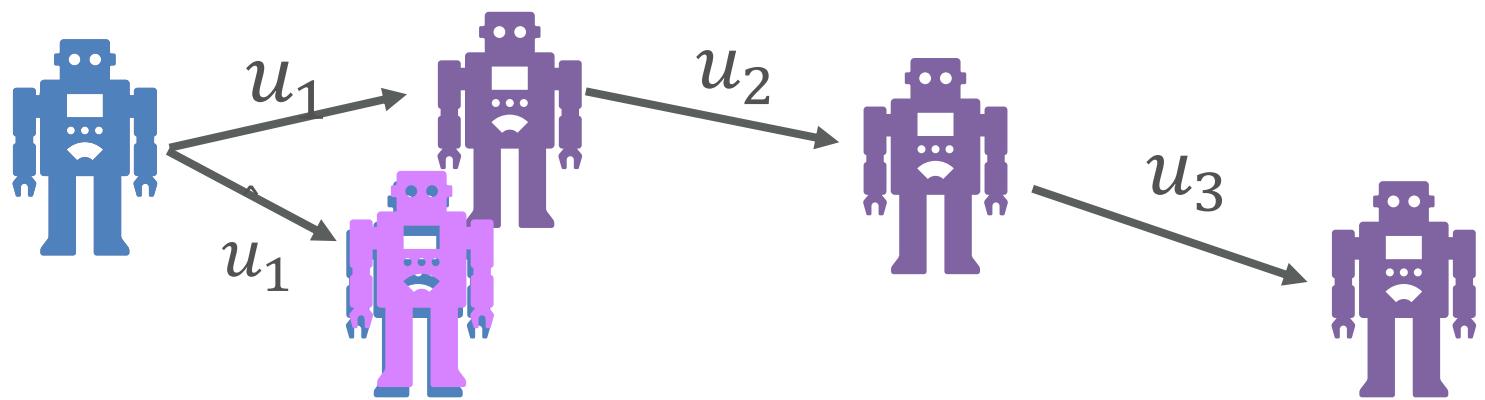


Execute Action

Plan

Ground Truth

MPC controller



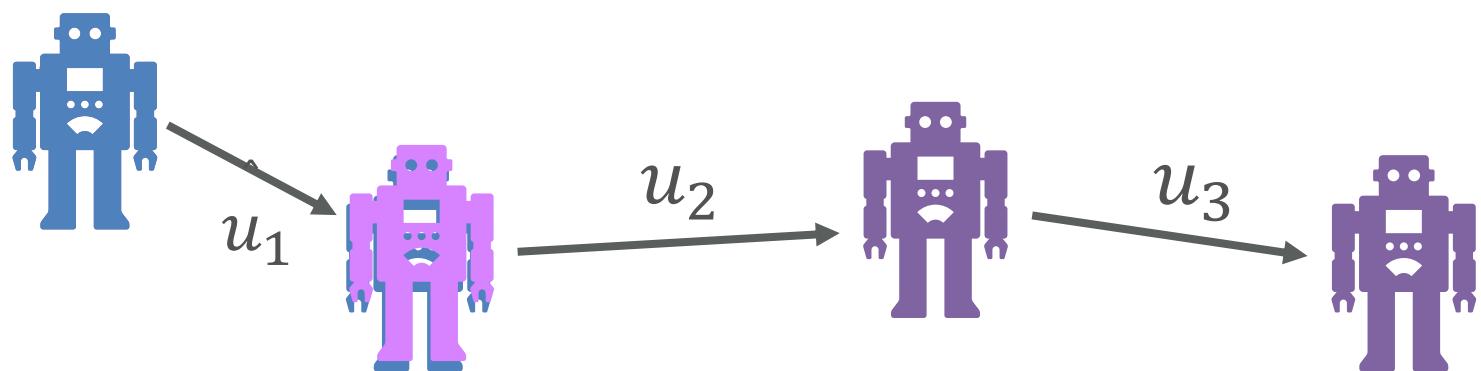
Estimate State

Plan

Ground Truth

State Estimate

MPC controller



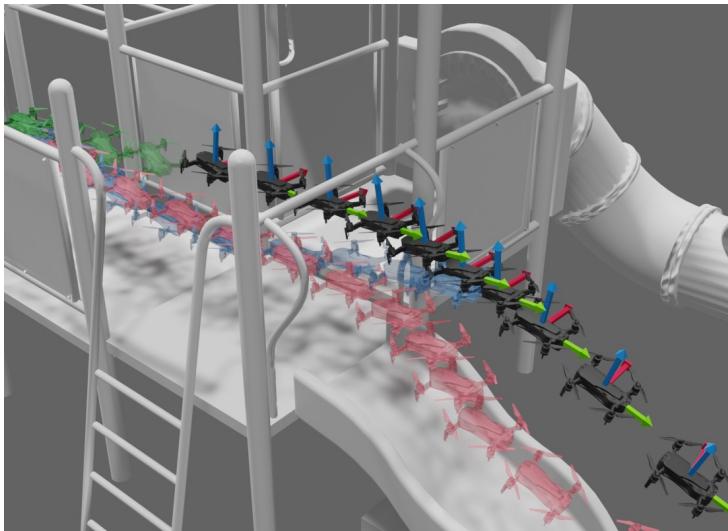
Replan

Plan

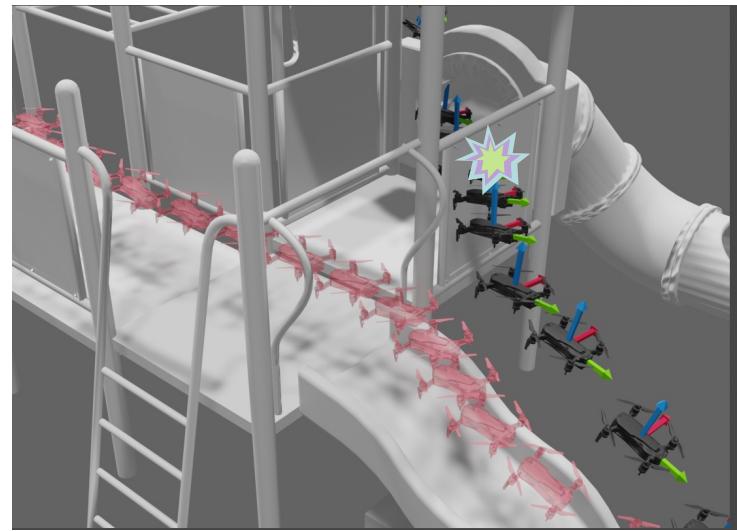
Ground Truth

State Estimate

MPC controller



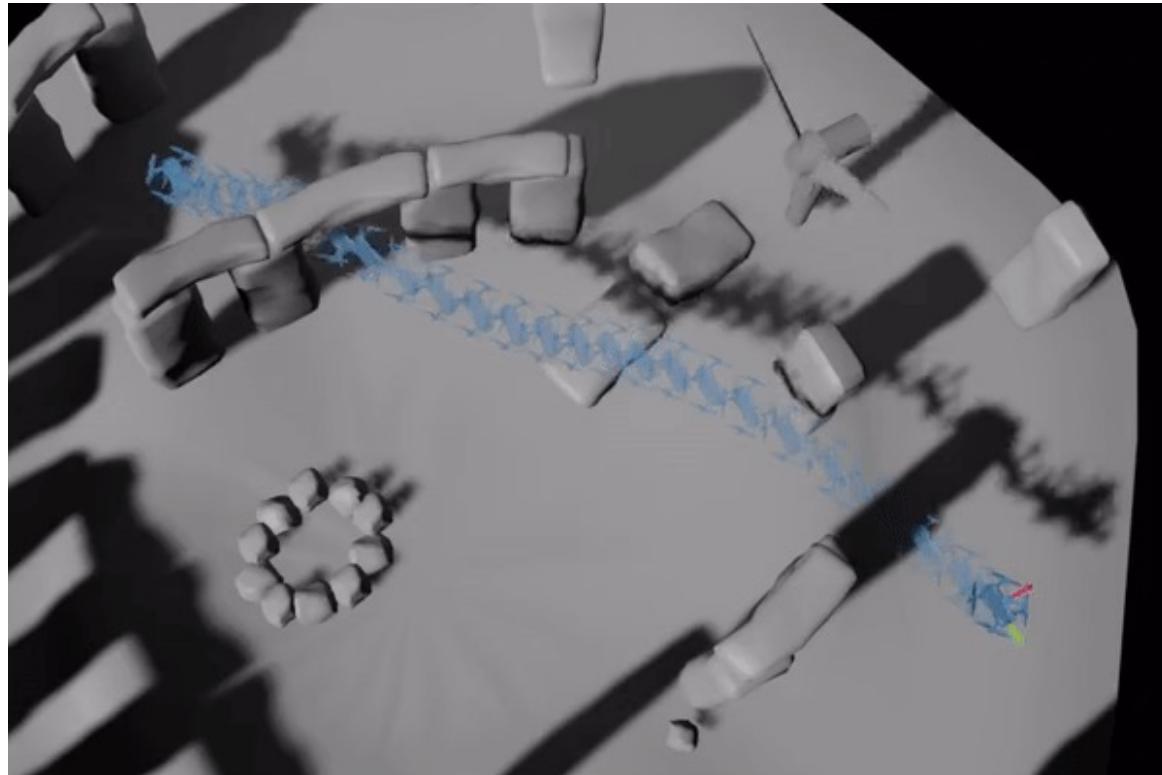
With replanning



Open loop

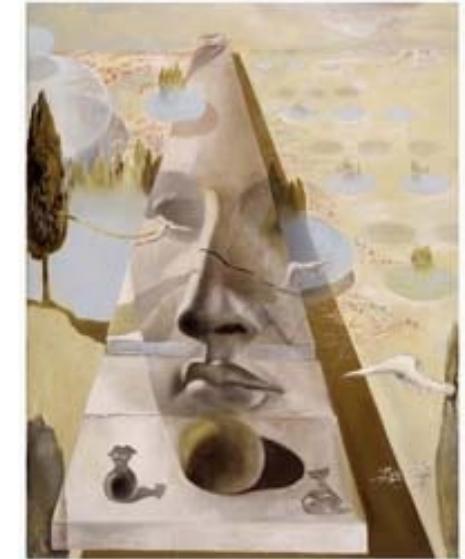
Red is initially planned trajectory

MPC controller



CS231A

Computer Vision: From 3D Reconstruction to Recognition



Next lecture:

Angjoo Kanazawa from UC Berkeley